

## TP - Monitoring des logs d'une application Spring Boot avec Filebeat, Logstash, Elasticsearch et Kibana

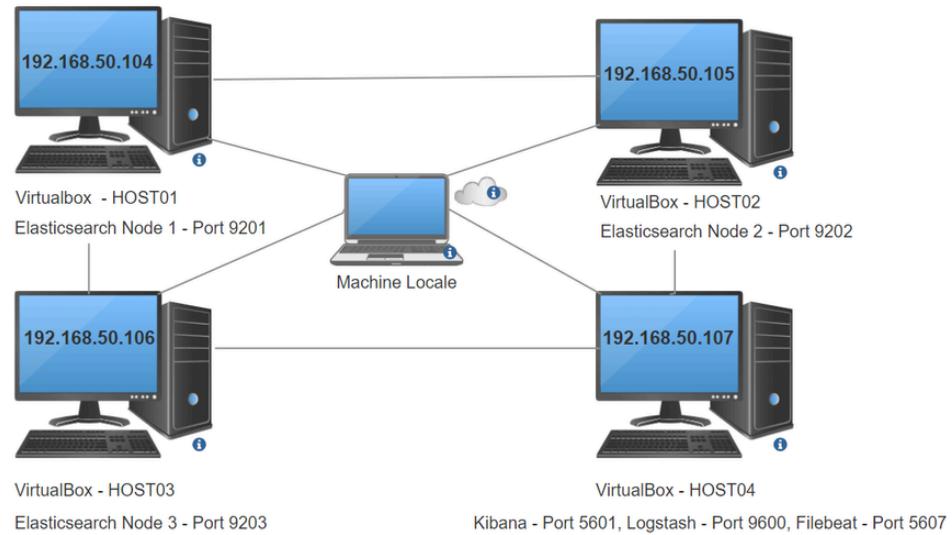
 Réalisé par **Mr. Iyanou Eraste AKANDE**, Ingénieur des données Telecom à Synaptique Maghreb, Ingénieur Certifié Elasticsearch.  
eraste.akande@gmail.com

### Etape 1: Introduction à la pile ELK

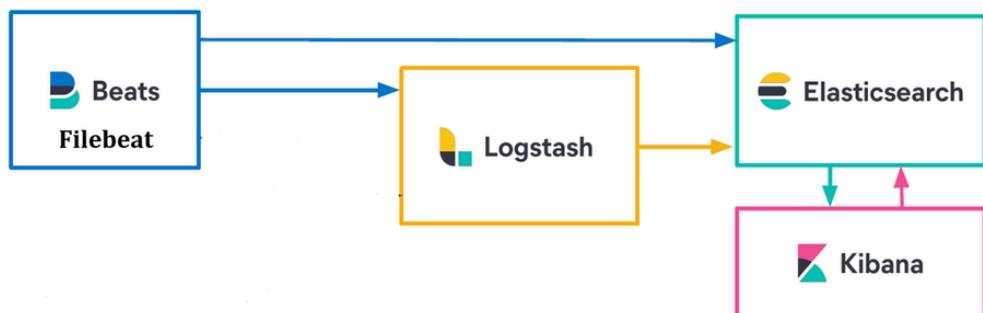
 Présentation des composantes de la pile ELK, le rôle de chaque composante ainsi que les différents cas d'utilisation → (PPT Introduction to Elastic Stack).

### Etape 2: Architecture

#### Les machines virtuelles



#### Les applications



## Etape 3: Installation et configuration des machines virtuelles

- Téléchargez et installez la version de Virtual Box compatible à votre système d'exploitation sur le site <https://www.virtualbox.org/wiki/Downloads>.
- Dans certains cas pour les utilisateurs Windows, vous devriez installer au préalable Visual Studio C++ Redistributable à partir du site <https://learn.microsoft.com/en-us/cpp/windows/latest-supported-vc-redist?view=msvc-170>.
- Téléchargez et installez la version de Vagrant compatible à votre système d'exploitation à partir du site [https://developer.hashicorp.com/vagrant/install?product\\_intent=vagrant](https://developer.hashicorp.com/vagrant/install?product_intent=vagrant).
- Téléchargez le fichier nommé Vagrantfile depuis ce répertoire GitHub <https://github.com/ianou/elastic-training-lab>.
- Placez le fichier Vagrantfile dans un dossier et exécutez ce qui suit depuis cet emplacement

```
1 cd $vagrant_file_repository
2 vagrant up
```

- Se connecter sur chaque machine et vérifier qu'on arrive à faire un ping vers les autres adresses IP

```
1 ping -c 10 HOST01
2 ping -c 10 HOST02
3 ping -c 10 HOST03
4 ping -c 10 HOST04
```

## Etape 4: Installation et configuration des applications

### 1. Les noeuds Elasticsearch

#### HOST01

- Se connecter au HOST01 depuis votre machine en utilisant Vagrant

```
1 cd $vagrant_file_repository
2 vagrant ssh HOST01
```

- Création du dossier d'installation

```
1 sudo mkdir -p /opt/training
```

- Téléchargement de Elasticsearch 8.13.0

```
1 cd /opt/training
2 sudo wget https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-8.13.0-linux-x86_64.tar.gz
```

- Décompression de Elasticsearch 8.13.0

```
1 sudo tar -xzf elasticsearch-8.13.0-linux-x86_64.tar.gz
2 sudo rm elasticsearch-8.13.0-linux-x86_64.tar.gz
```

- Modifier le fichier de configuration `elasticsearch.yml` avec le contenu du fichier `es_node01.yml` situé dans le répertoire GitHub <https://github.com/ianou/elastic-training-lab>.

```
1 cd /opt/training/elasticsearch-8.13.0
2 sudo rm config/elasticsearch.yml
```

```
3 sudo nano config/elasticsearch.yml
```

- Mettre à jour le paramètre `vm.max_map_count`

```
1 sudo grep vm.max_map_count /etc/sysctl.conf
2
3 grep 'vm.max_map_count=262144' /etc/sysctl.conf || sudo echo 'vm.max_map_count=262144'
4 | sudo tee -a /etc/sysctl.conf
5
6 sudo sysctl -p
```

- Donner les permissions nécessaires à l'utilisateur sur le dossier `elasticsearch`

```
1 sudo chown -R vagrant:vagrant /opt/training/elasticsearch-8.13.0
```

- Générer les certificats de connexion Elasticsearch en utilisant les fichiers contenus dans le répertoire GitHub

<https://github.com/ianou/elastic-training-lab/certificates>.

```
1 sudo mkdir -p config/certificates/ca
2 sudo mkdir -p config/certificates/node01
3 sudo nano config/certificates/node01/node01.key
4 sudo nano config/certificates/node01/node01.crt
5 sudo nano config/certificates/ca/ca.crt
```

- 💡 Voici comment tous les certificats ont été générés en utilisant le fichier `instances.yml` sur le répertoire GitHub <https://github.com/ianou/elastic-training-lab/>.

```
1 cd /opt/training/elasticsearch-8.13.0
2 sudo mkdir config/certificates
3 sudo nano config/certificates/instances.yml
4 sudo bin/elasticsearch-certutil ca --silent -- pem -out config/certificates/ca.zip
5 sudo apt install unzip
6 sudo unzip -o config/certificates/ca.zip -d config/certificates
7
8 sudo bin/elasticsearch-certutil cert --silent -- pem -out config/certificates/bundle.zip
9 --in config/certificates/instances.yml --ca-cert config/certificates/ca/ca.crt --ca-key
10 config/certificates/ca/ca.key
11
12 sudo unzip -o config/certificates/bundle.zip -d config/certificates
13 sudo rm config/certificates/bundle.zip
14 sudo rm config/certificates/ca.zip
```

- Autoriser les ports sur le pare-feu

```
1 sudo ufw enable
2 sudo ufw allow 9201/tcp
3 sudo ufw allow 9301/tcp
4 sudo ufw status
```

- Démarrer le nœud Elasticsearch comme un daemon et vérifier les logs

```
1 ./bin/elasticsearch -d -p pid
```

- 💡 Pour arrêter Elasticsearch, exécutez `pskill -F pid`.

## HOST02 ☁

- Se connecter au HOST02 depuis votre machine en utilisant Vagrant

```
1 cd $vagrant_file_repository
2 vagrant ssh HOST02
```

- Création du dossier d'installation

```
1 sudo mkdir -p /opt/training
```

- Téléchargement de Elasticsearch 8.13.0

```
1 cd /opt/training
2 sudo wget https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-8.13.0-linux-x86_64.tar.gz
```

- Décompression de Elasticsearch 8.13.0

```
1 sudo tar -xzf elasticsearch-8.13.0-linux-x86_64.tar.gz
2 sudo rm elasticsearch-8.13.0-linux-x86_64.tar.gz
```

- Modifier le fichier de configuration `elasticsearch.yml` avec le contenu du fichier `es_node02.yml` situé dans le répertoire GitHub <https://github.com/yanou/elastic-training-lab>.

```
1 cd /opt/training/elasticsearch-8.13.0
2 sudo rm config/elasticsearch.yml
3 sudo nano config/elasticsearch.yml
```

- Mettre à jour le paramètre `vm.max_map_count`

```
1 sudo grep vm.max_map_count /etc/sysctl.conf
2
3 grep 'vm.max_map_count=262144' /etc/sysctl.conf || sudo echo 'vm.max_map_count=262144'
4 | sudo tee -a /etc/sysctl.conf
5
6 sudo sysctl -p
```

- Donner les permissions nécessaires à l'utilisateur sur le dossier `elasticsearch`

```
1 sudo chown -R vagrant:vagrant /opt/training/elasticsearch-8.13.0
```

- Générer les certificats de connexion Elasticsearch en utilisant les fichiers contenus dans le répertoire GitHub <https://github.com/yanou/elastic-training-lab/certificates> .

```
1 sudo mkdir -p config/certificates/ca
2 sudo mkdir -p config/certificates/node02
3 sudo nano config/certificates/node02/node02.key
4 sudo nano config/certificates/node02/node02.crt
5 sudo nano config/certificates/ca/ca.crt
```

- Autoriser les ports sur le pare-feu

```
1 sudo ufw enable
2 sudo ufw allow 9202/tcp
3 sudo ufw allow 9302/tcp
4 sudo ufw status
```

- Démarrer le nœud Elasticsearch comme un daemon et vérifier les logs

```
1 ./bin/elasticsearch -d -p pid
```

## HOST03 ↗

- Se connecter au HOST03 depuis votre machine en utilisant Vagrant

```
1 cd $vagrant_file_repository  
2 vagrant ssh HOST03
```

- Création du dossier d'installation

```
1 sudo mkdir -p /opt/training
```

- Téléchargement de Elasticsearch 8.13.0

```
1 cd /opt/training  
2 sudo wget https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-8.13.0-linux-x86_64.tar.gz
```

- Décompression de Elasticsearch 8.13.0

```
1 sudo tar -xzf elasticsearch-8.13.0-linux-x86_64.tar.gz  
2 sudo rm elasticsearch-8.13.0-linux-x86_64.tar.gz
```

- Modifier le fichier de configuration `elasticsearch.yml` avec le contenu du fichier `es_node03.yml` situé dans le répertoire GitHub <https://github.com/yanou/elastic-training-lab>.

```
1 cd /opt/training/elasticsearch-8.13.0  
2 sudo rm config/elasticsearch.yml  
3 sudo nano config/elasticsearch.yml
```

- Mettre à jour le paramètre `vm.max_map_count`

```
1 sudo grep vm.max_map_count /etc/sysctl.conf  
2  
3 grep 'vm.max_map_count=262144' /etc/sysctl.conf || sudo echo 'vm.max_map_count=262144'  
4 | sudo tee -a /etc/sysctl.conf  
5  
6 sudo sysctl -p
```

- Donner les permissions nécessaires à l'utilisateur sur le dossier `elasticsearch`

```
1 sudo chown -R vagrant:vagrant /opt/training/elasticsearch-8.13.0
```

- Générer les certificats de connexion Elasticsearch en utilisant les fichiers contenus dans le répertoire GitHub <https://github.com/yanou/elastic-training-lab/certificates> .

```
1 sudo mkdir -p config/certificates/ca  
2 sudo mkdir -p config/certificates/node03  
3 sudo nano config/certificates/node03/node03.key  
4 sudo nano config/certificates/node03/node03.crt  
5 sudo nano config/certificates/ca/ca.crt
```

- Autoriser les ports sur le pare-feu

```
1 sudo ufw enable  
2 sudo ufw allow 9203/tcp  
3 sudo ufw allow 9303/tcp  
4 sudo ufw status
```

- Démarrer le nœud Elasticsearch comme un daemon et vérifier les logs

```
1 ./bin/elasticsearch -d -p pid
```

## Connexion au Cluster Elasticsearch 🔒

- Générer le mot de passe de l'utilisateur `elastic` et `kibana_system` en utilisant l'un des hosts et sauvegarder ces mots de passe

```
1 cd /opt/training/elasticsearch-8.13.0  
2 bin/elasticsearch-reset-password -u elastic  
3 bin/elasticsearch-reset-password -u kibana_system
```

- Sur votre machine local, aller sur votre navigateur et accéder au lien [https://192.168.50.104:9201/\\_cluster/health?pretty](https://192.168.50.104:9201/_cluster/health?pretty) pour vérifier que le cluster est en bon état (statut "green").
- Sur votre machine local, aller sur votre navigateur et accéder au lien [https://192.168.50.104:9201/\\_cat/nodes?pretty](https://192.168.50.104:9201/_cat/nodes?pretty) pour vérifier que les 3 nœuds ont rejoint le cluster.

- Au lieu d'aller sur le navigateur, on peut faire les mêmes vérifications depuis l'un des hosts avec les commandes suivantes :

```
1 curl --cacert /opt/training/elasticsearch-8.13.0/config/certificates/ca/ca.crt  
2 -u elastic:$elastic_password -XGET 'https://192.168.50.104:9201/_cluster/health?pretty'  
3  
4 curl --cacert /opt/training/elasticsearch-8.13.0/config/certificates/ca/ca.crt  
5 -u elastic:$elastic_password -XGET 'https://192.168.50.104:9201/_cat/nodes?pretty'
```

## 2. Le serveur Kibana

- La configuration de Kibana se fera sur le HOST04

- Se connecter au HOST04 depuis votre machine en utilisant Vagrant

```
1 cd $vagrant_file_repository  
2 vagrant ssh HOST04
```

- Création du dossier d'installation

```
1 sudo mkdir -p /opt/training
```

- Téléchargement de Kibana 8.13.0

```
1 cd /opt/training  
2 sudo wget https://artifacts.elastic.co/downloads/kibana/kibana-8.13.0-amd64.deb
```

- Décompression de Kibana 8.13.0

```
1 sudo dpkg -i kibana-8.13.0-amd64.deb  
2 sudo rm kibana-8.13.0-amd64.deb
```

- Modifier le fichier de configuration `kibana.yml` avec le contenu du fichier `kibana.yml` situé dans le répertoire GitHub <https://github.com/ianou/elastic-training-lab>. Ne pas oublier de mettre à jour le mot de passe du user `kibana_system`.

```
1 sudo rm /etc/kibana/kibana.yml  
2 sudo nano /etc/kibana/kibana.yml
```

- Déposer le certificat d'autorité de Elasticsearch sur le HOST04. Copier le certificat depuis le répertoire GitHub <https://github.com/ianou/elastic-training-lab/certificates>.

```
1 sudo mkdir -p /opt/training/ca  
2 sudo nano /opt/training/ca/ca.crt
```

- Autoriser les ports sur le pare-feu

```
1 sudo ufw enable  
2 sudo ufw allow 5601/tcp  
3 sudo ufw status
```

- Démarrer Kibana

```
1 sudo /bin/systemctl daemon-reload  
2 sudo systemctl enable kibana.service  
3 sudo systemctl start kibana.service  
4 sudo systemctl status kibana.service
```

 Pour arrêter le service Kibana, exécutez `sudo systemctl stop kibana.service`.

- Afficher les logs du service Kibana

```
1 sudo journalctl --unit=kibana.service -n 100 --no-pager
```

- Sur votre machine aller sur l'interface de Kibana <http://192.168.50.107:5601> et connecter vous avec le user `elastic` et son mot de passe généré précédemment.

### 3. L'ETL Logstash

 La configuration de Logstash se fera sur le HOST04

- Se connecter au HOST04 depuis votre machine en utilisant Vagrant

```
1 cd $vagrant_file_repository  
2 vagrant ssh HOST04
```

- Téléchargement de Logstash 8.13.0

```
1 cd /opt/training  
2 sudo wget https://artifacts.elastic.co/downloads/logstash/logstash-8.13.0-amd64.deb
```

- Décompression de Logstash 8.13.0

```
1 sudo dpkg -i logstash-8.13.0-amd64.deb  
2 sudo rm logstash-8.13.0-amd64.deb
```

- Modifier le fichier de configuration `logstash.yml` avec le contenu du fichier `logstash.yml` situé dans le répertoire GitHub <https://github.com/ilyanou/elastic-training-lab>.

```
1 sudo rm /etc/logstash/logstash.yml  
2 sudo nano /etc/logstash/logstash.yml
```

- Autoriser les ports sur le pare-feu

```
1 sudo ufw enable  
2 sudo ufw allow 9600/tcp  
3 sudo ufw allow 5044/tcp  
4 sudo ufw status
```

- Donner les permissions nécessaires à Logstash

```
1 sudo chmod 755 /usr/share/logstash/data  
2 sudo chown -R logstash:logstash /usr/share/logstash/data
```

- Démarrer Logstash

```
1 sudo /bin/systemctl daemon-reload
2 sudo systemctl enable logstash.service
3 sudo systemctl start logstash.service
4 sudo systemctl status logstash.service
```

**i** Pour arrêter le service Logstash, exécuter `sudo systemctl stop logstash.service`.

- Afficher les logs du service Logstash

```
1 sudo journalctl --unit=logstash.service -n 100 --no-pager
```

- Aller sur l'interface de Logstash sur votre navigateur avec le lien <http://192.168.50.107:9600/?pretty> pour se rassurer que Logstash est démarré.

#### 4. L'agent Filebeat

**i** La configuration de Filebeat se fera sur le HOST04

- Se connecter au HOST04 depuis votre machine en utilisant Vagrant

```
1 cd $vagrant_file_repository
2 vagrant ssh HOST04
```

- Téléchargement de Filebeat 8.13.0

```
1 cd /opt/training
2 sudo wget https://artifacts.elastic.co/downloads/beats/filebeat/filebeat-8.13.0-amd64.deb
```

- Décompression de Filebeat 8.13.0

```
1 sudo dpkg -i filebeat-8.13.0-amd64.deb
2 sudo rm filebeat-8.13.0-amd64.deb
```

- Autoriser les ports sur le pare-feu

```
1 sudo ufw enable
2 sudo ufw allow 5067/tcp
3 sudo ufw status
```

- Démarrer Filebeat

```
1 sudo /bin/systemctl daemon-reload
2 sudo systemctl enable filebeat.service
3 sudo systemctl start filebeat.service
4 sudo systemctl status filebeat.service
```

**i** Pour arrêter le service Filebeat, exécuter `sudo systemctl stop logstash.service`.

- Afficher les logs du service Filebeat

```
1 sudo journalctl --unit=filebeat.service -n 100 --no-pager
```

## Etape 5: Introduction à Filebeat ↗

- 💡 Explication du fonctionnement de Filebeat et présentation des différentes options et paramètres de configuration → (PPT Starting with Filebeat).

## Etape 6: Introduction à Logstash ↗

- 💡 Explication du fonctionnement de Logstash et présentation des différentes options et paramètres de configuration des pipelines Logstash → (PPT Starting with Logstash).

## Etape 7: Introduction à Elasticsearch (Mapping et Analysers) ↗

- 💡 Explication de quelques concepts importants liés à Elasticsearch et à son fonctionnement → (PPT Starting with Elasticsearch).

## Etape 8: Collecte des logs Spring Boot : De Filebeat vers Logstash et transfert vers Elasticsearch ↗

- 💡 Ce LAB se déroulera sur le HOST04.

- Se connecter au HOST04 depuis votre machine en utilisant Vagrant

```
1 cd $vagrant_file_repository
2 vagrant ssh HOST04
```

- Créer le fichier `spring-boot-app-logs.json` en utilisant le fichier `spring-boot-app-logs.json` situé dans le répertoire GitHub <https://github.com/ianou/elastic-training-lab>.

```
1 sudo mkdir -p /opt/training/app/logs
2 sudo nano /opt/training/app/logs/spring-boot-app-logs.json
```

- Créer le pipeline Logstash pour le traitement des données venant de Filebeat. Utiliser le fichier `filebeat-pipeline.conf` situé dans le répertoire GitHub <https://github.com/ianou/elastic-training-lab>.

```
1 sudo nano /etc/logstash/conf.d/filebeat-pipeline.conf
```

- Redémarrer le service Logstash

```
1 sudo systemctl restart logstash.service
```

- Mettre à jour le index template de filebeat en utilisant le fichier `filebeat-index-template.txt` situé dans le répertoire GitHub <https://github.com/ianou/elastic-training-lab>.
- Configuration de `filebeat.yml` en utilisant le fichier `filebeat.yml` situé dans le répertoire GitHub <https://github.com/ianou/elastic-training-lab>.

```
1 sudo rm /etc/filebeat/filebeat.yml
```

```
2 sudo nano /etc/filebeat/filebeat.yml
```

- Supprimer le registre de Filebeat

```
1 sudo rm -r /var/lib/filebeat/registry  
2 sudo rm /var/lib/filebeat/meta.json  
3 sudo rm /var/lib/filebeat/filebeat.lock
```

- Redémarrer le service Filebeat

```
1 sudo systemctl restart filebeat.service
```

- Afficher les logs du service Filebeat

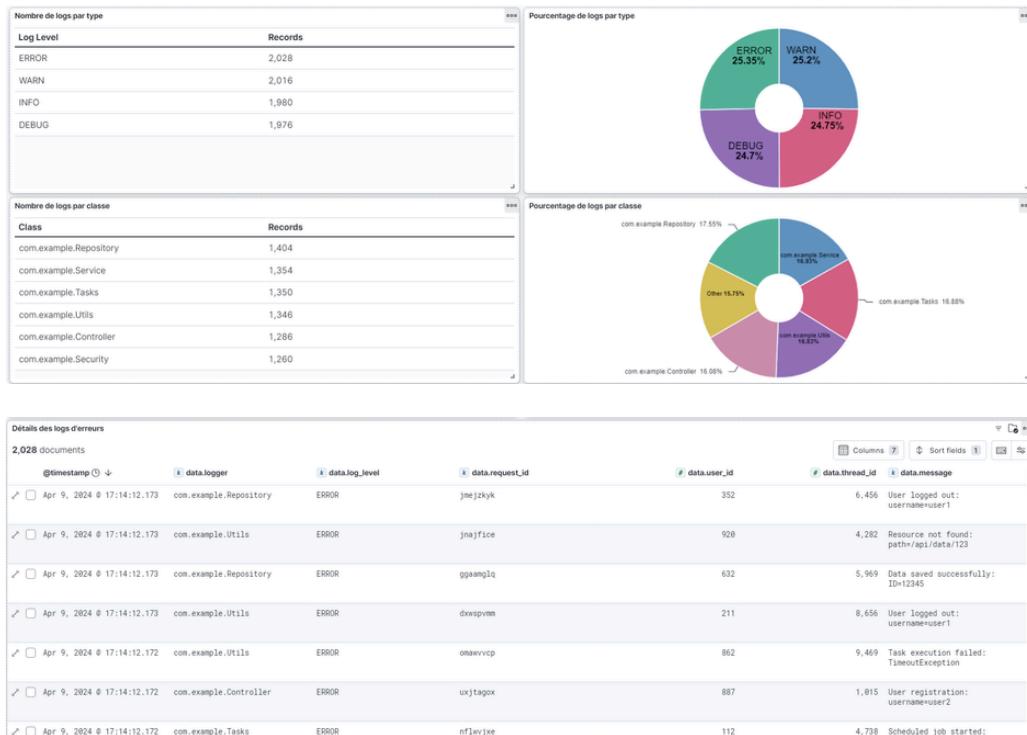
```
1 sudo journalctl --unit=filebeat.service -n 100 --no-pager
```

- Afficher les logs du service Logstash

```
1 sudo journalctl --unit=logstash.service -n 100 --no-pager
```

- Aller sur Kibana, choisir le data view nommé “**filebeat**” et aller dans la rubrique “**Discover**” pour vérifier que les logs ont été envoyés sur le cluster Elasticsearch.
- Tester sur Discover que l’analyser `filebeat_analyser` fonctionne comme voulu

## Etape 9: Visualisation des Logs et création de Dashboards avec Kibana ↗



## Tester que le système est automatique

Sur le HOST04,

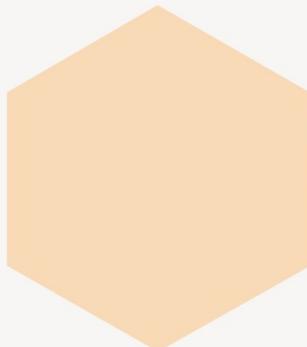
- Créer un nouveau fichier `spring-boot-app-logs-test.json` en utilisant le fichier `spring-boot-app-logs-test.json` situé dans le répertoire GitHub <https://github.com/yanou/elastic-training-lab>.

```
1 sudo nano /opt/training/app/logs/spring-boot-app-logs-test.json
```

- Visualiser que les logs ont mis le Dashboard à jour.

# Introduction to Elastic Stack

Mr. Iyanou Eraste AKANDE  
Elastic Certified Engineer  
Data Engineer at Synaptique Maghreb

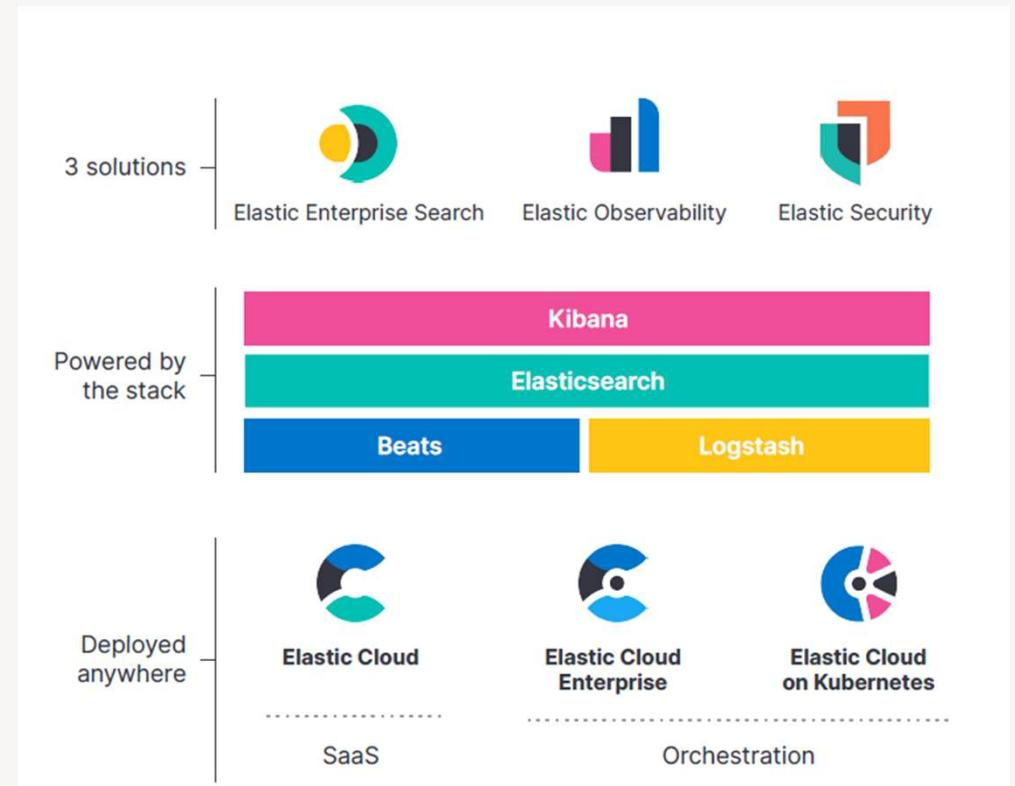


# The Beginning

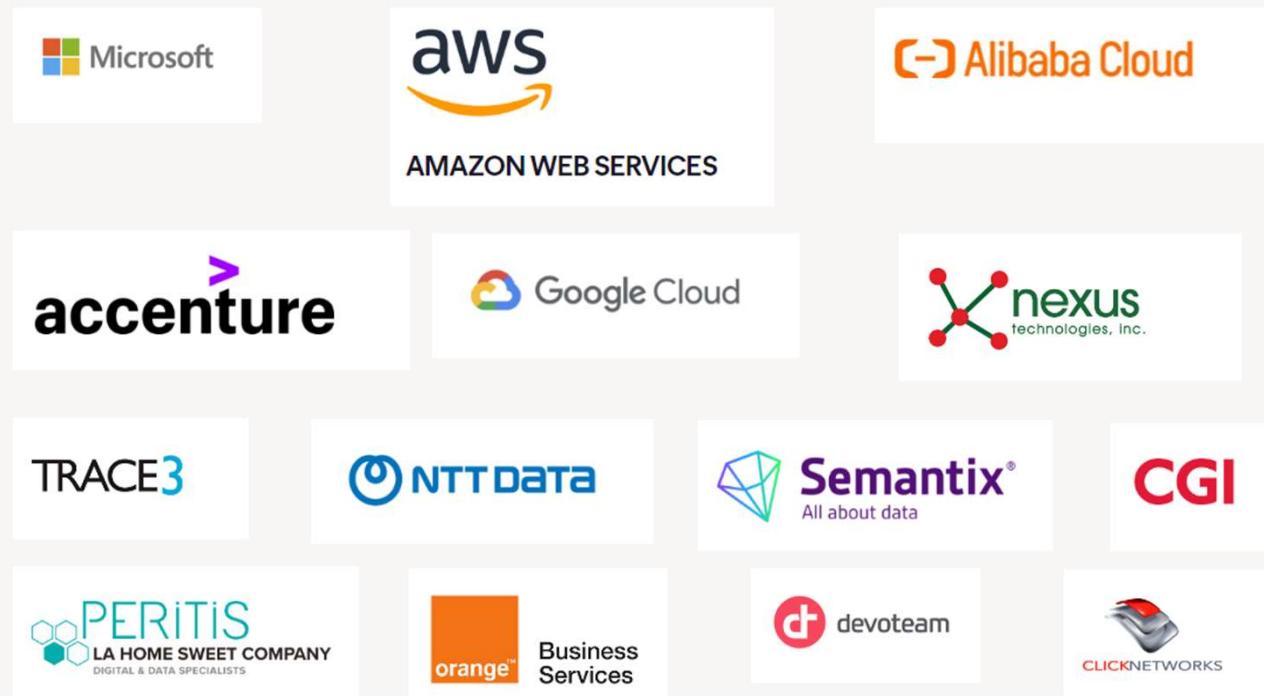
Lucene	Compass	Elasticsearch
1999	2004	2010
Doug Cutting	Shay Banon	Shay Banon
Search Engine in Java	Search Product in Java, scalability included, built on Lucene	Search product built on Compass concepts, distributed and integrated with other language

# Elastic Stack Components

- Elasticsearch : NoSQL Database
- Kibana : Analytics and Visualization platform
- Logstash : Server side data processing pipeline
- Beats : Data Shippers (Filebeat, Metricbeat, Heartbeat, Elastic Agents)



# Elastic Partners



<https://partners.elastic.co/findapartner/>

Introduction to Elastic Stack

# Elastic Stack Use Cases



- **Log Management and Analysis:** One of the primary use cases for Elastic Stack is log management and analysis. Organizations use it to collect, parse, index, and analyze logs generated by various systems, applications, and services.
- **Security Information and Event Management (SIEM):** Elastic Stack can be used as a SIEM solution to collect, analyze, and visualize security-related data such as logs, network traffic, and system events.
- **Real-time Monitoring and Alerting:** Elastic Stack can be utilized for real-time monitoring of infrastructure, applications, and services. By collecting metrics, logs, and events in real-time, organizations can set up alerts and notifications.
- **Application Performance Monitoring (APM):** Elastic APM, an extension of the Elastic Stack, is used for monitoring the performance of applications and microservices.

# Elastic Stack Use Cases



- **Business Analytics and Insights:** Elastic Stack can be leveraged for business analytics and data visualization purposes. By indexing and analyzing large volumes of data from different sources, organizations can gain valuable insights into customer behavior, market trends, and operational performance.
- **Search and Information Retrieval:** Elasticsearch supports full-text search, fuzzy search, geospatial search, and faceted navigation, making it suitable for building search-driven applications and platforms.
- **Data Exploration and Discovery:** Elastic Stack can be used for data exploration and discovery in various domains such as scientific research, e-commerce, and content management.

# Elasticsearch

NoSQL Database



first_name	name	gender	city	country
Ali	KOZNI	M	Rabat	Morocco
Achille	BADRE	F	Tetouan	Morocco

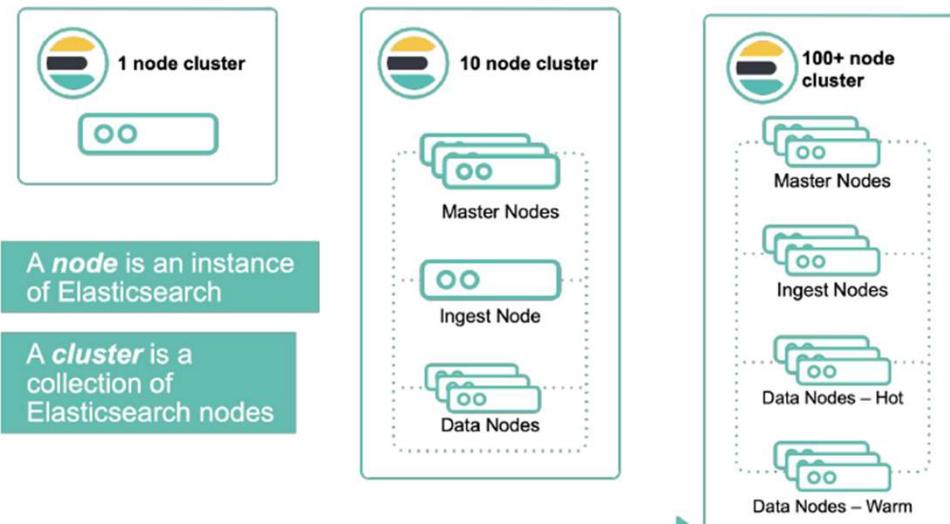


```
{  
  "first_name" : "Ali",  
  "name" : "KOZNI",  
  "gender" : "M",  
  "city" : "Rabat",  
  "country" : "Morocco"  
},  
{  
  "first_name" : "Achille",  
  "name" : "BADRE",  
  "gender" : "F",  
  "city" : "Tetouan",  
  "country" : "Morocco"  
}
```

Introduction to Elastic Stack

# Elasticsearch

Distributed and Horizontally Scalable

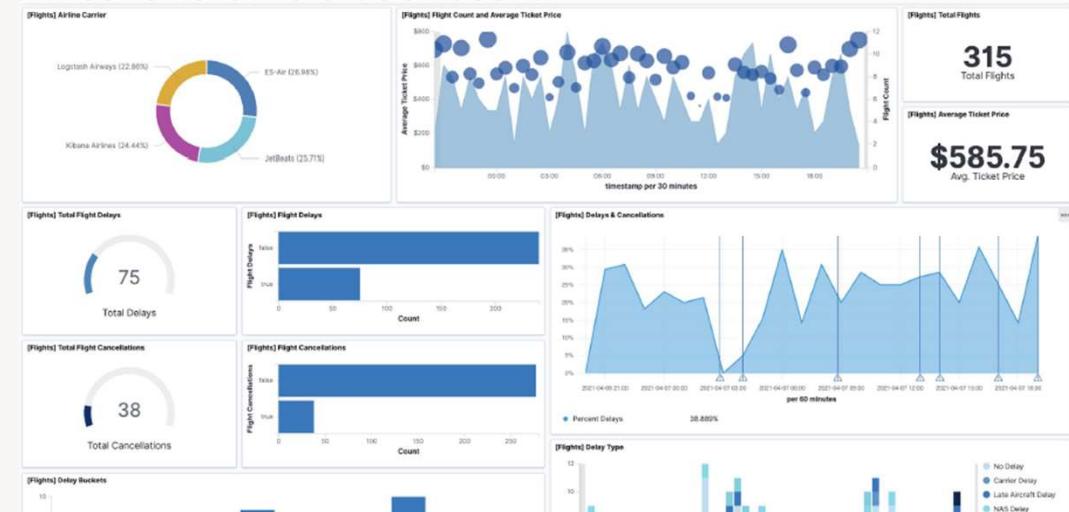


Introduction to Elastic Stack

# Kibana



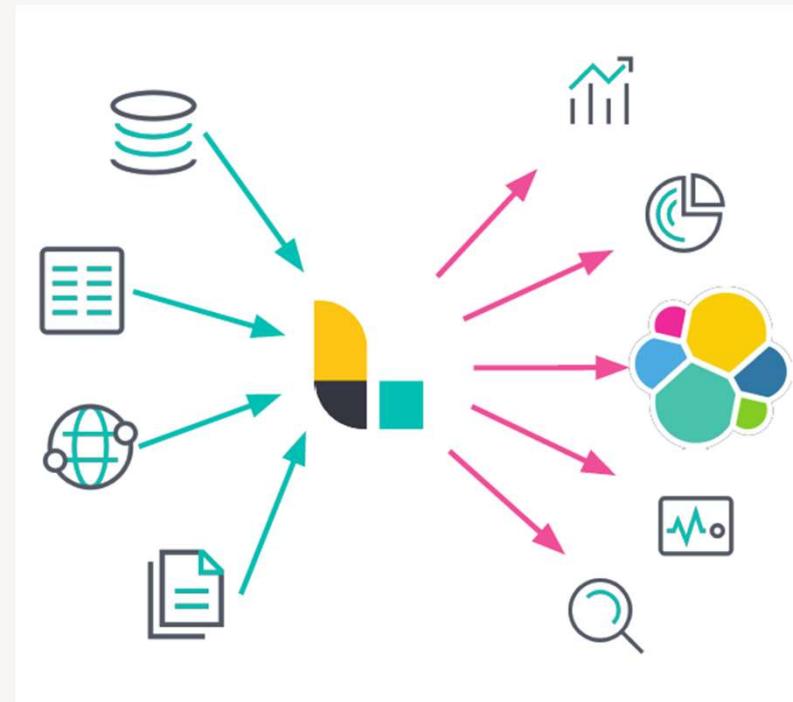
Analytics and Visualization Platform



Introduction to Elastic Stack

# Logstash

Data Processing Pipeline



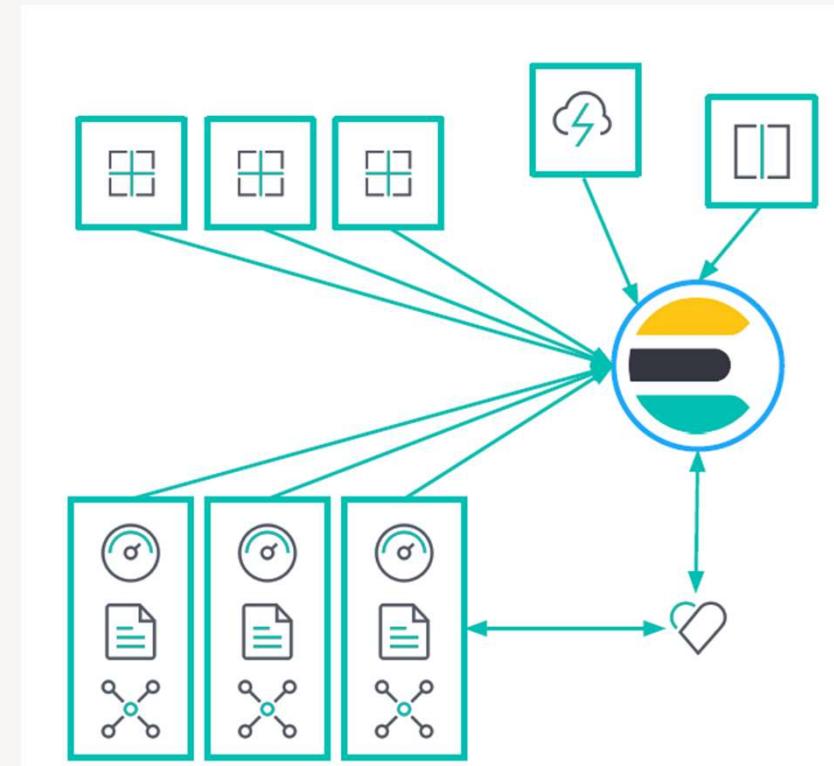
Introduction to Elastic Stack

# Beats



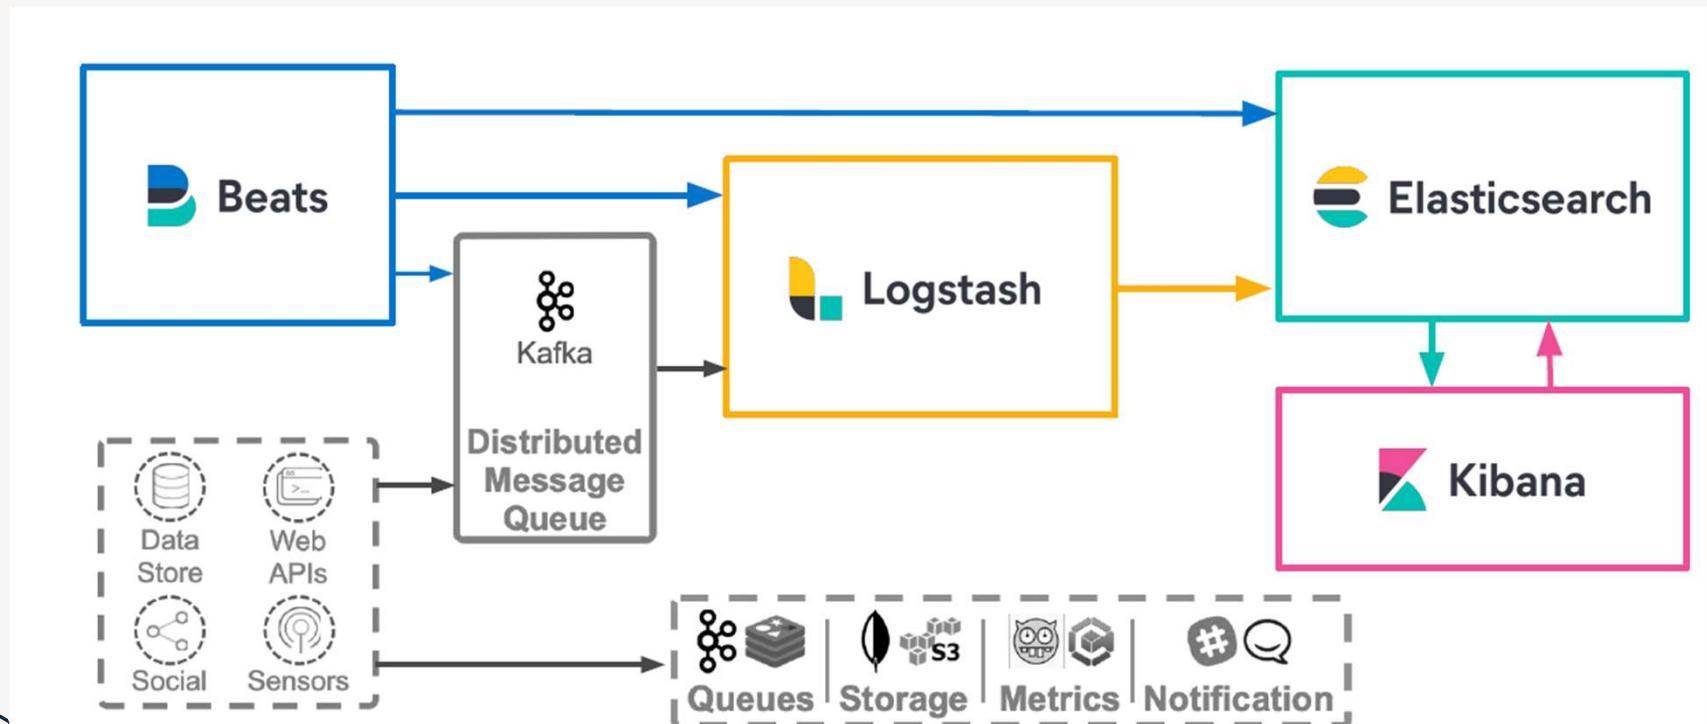
Open Source Data Shippers

Filebeat, Metricbeat, Heartbeat,  
Packetbeat, Winlogbeat, etc.



Introduction to Elastic Stack

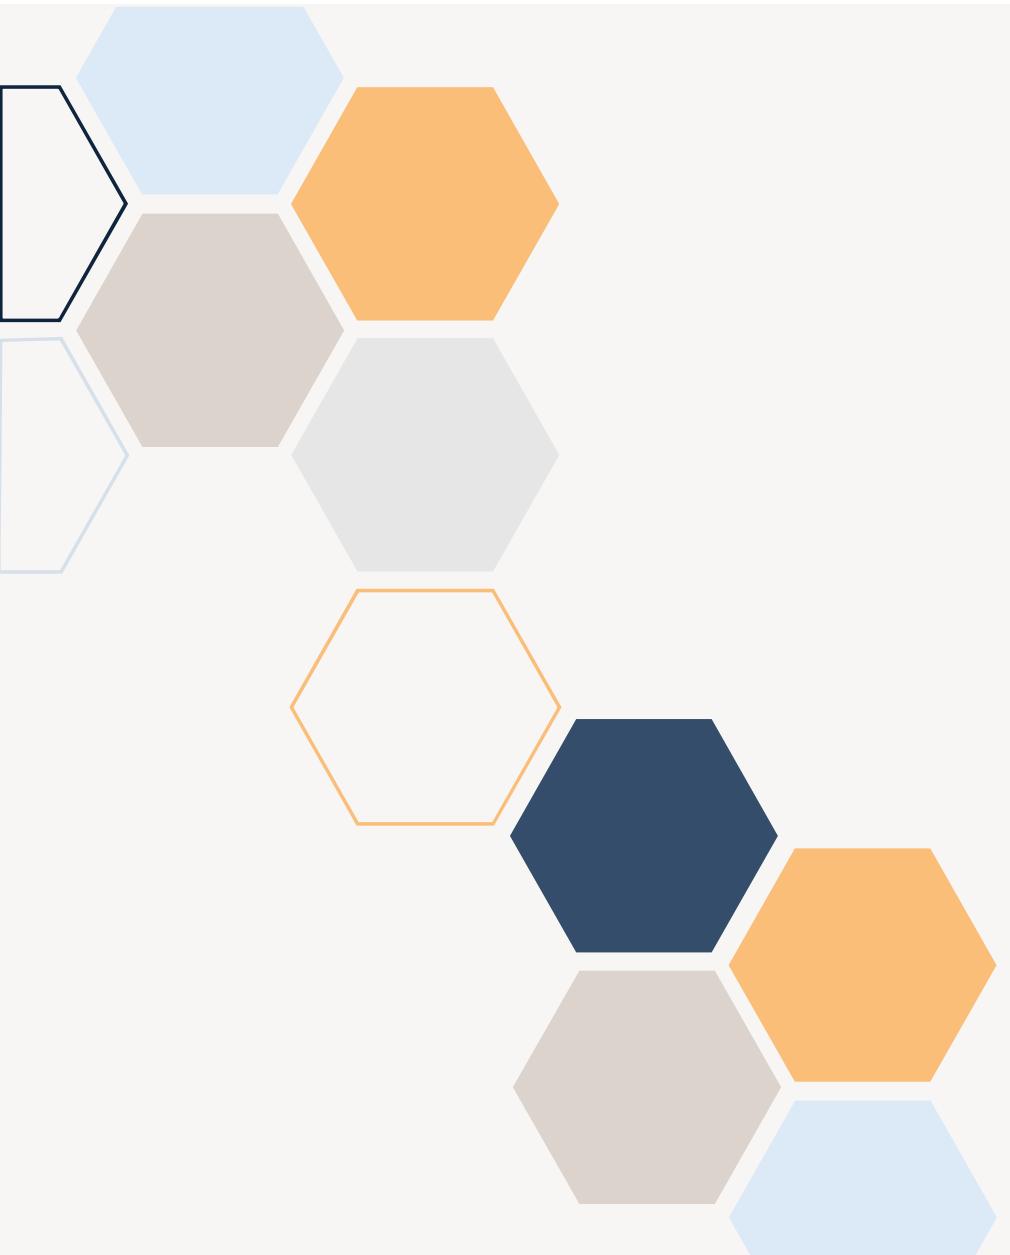
# Elastic Stack Architecture



Introduction to Elastic Stack

# Questions





# Thank you

Mr. Iyanou Eraste AKANDE  
[eraste.akande@gmail.com](mailto:eraste.akande@gmail.com)

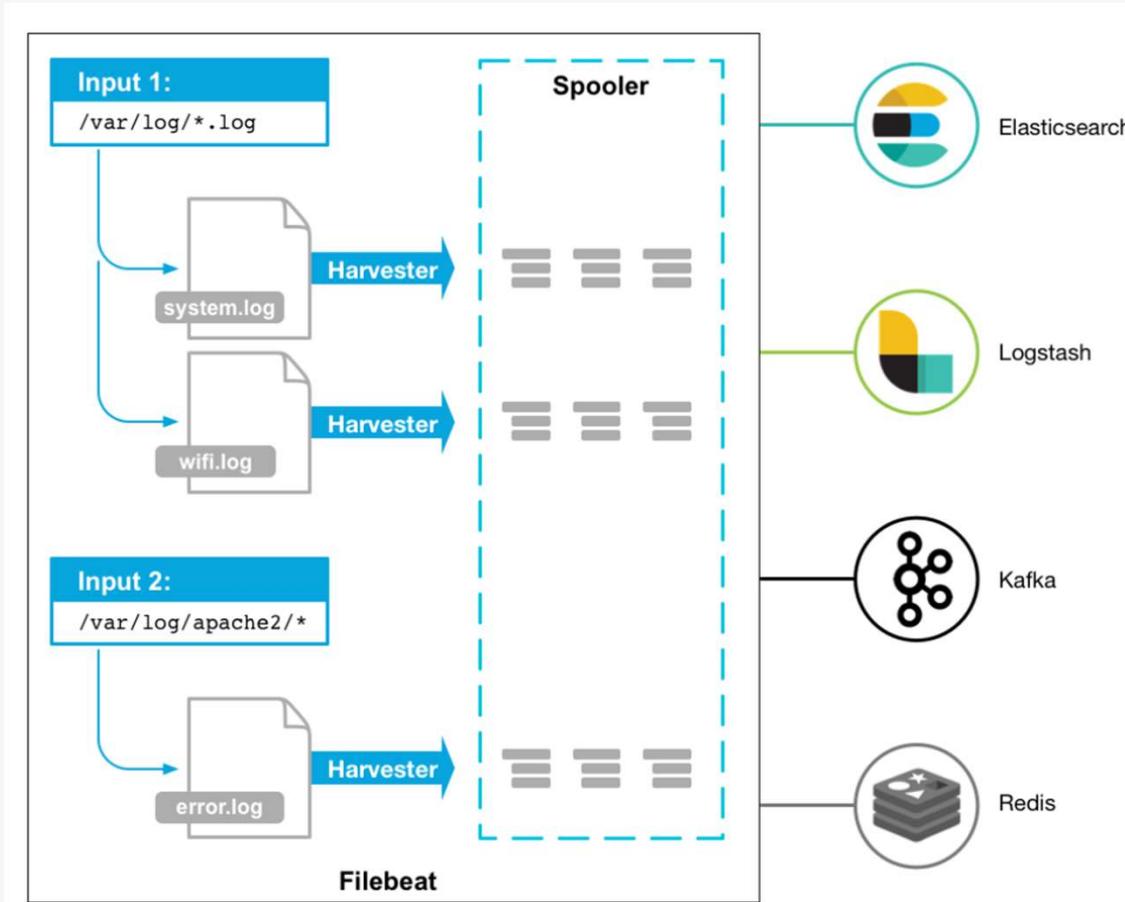
# Starting with Filebeat

Mr. Iyanou Eraste AKANDE  
Elastic Certified Engineer  
Data Engineer at Synaptique Maghreb



# Filebeat

## Architecture



# Filebeat



## Installation

- Download Filebeat and decompress it
  - [https://artifacts.elastic.co/downloads/beats/filebeat/filebeat-8.12.2-windows-x86\\_64.zip](https://artifacts.elastic.co/downloads/beats/filebeat/filebeat-8.12.2-windows-x86_64.zip)
  - `curl -L -O https://artifacts.elastic.co/downloads/beats/filebeat/filebeat-8.12.2-linux-x86_64.tar.gz && tar xzvf filebeat-8.12.2-linux-x86_64.tar.gz`
- Setup Filebeat (`cd filebeat_directory && filebeat setup -e`)
- Start Filebeat (`cd filebeat_directory && filebeat -e`)

# Filebeat



Filebeat.yml >> Input

## Input Types

- AWS CloudWatch
- AWS S3
- Azure Event Hub
- Azure Blob Storage
- CEL
- Cloud Foundry
- CometD
- Container
- Entity Analytics
- filestream
- GCP Pub/Sub
- HTTP Endpoint
- HTTP JSON
- journald
- Kafka
- Log (deprecated in 7.16.0, use filestream)
- MQTT
- NetFlow
- Office 365 Management Activity API
- Redis
- Stdin
- Syslog
- TCP
- UDP
- Google Cloud Storage

# Filebeat



Filebeat.yml >> Input

## Filestream Input

<https://www.elastic.co/guide/en/beats/filebeat/current/filebeat-input-filestream.html>

```
filebeat.inputs:
- type: filestream
  id: my-filestream-id
  paths:
    - /var/log/messages
    - /var/log/*.log
```

```
filebeat.inputs:
- type: filestream ❶
  id: my-filestream-id
  paths:
    - /var/log/system.log
    - /var/log/wifi.log
- type: filestream ❷
  id: apache-filestream-id
  paths:
    - "/var/log/apache2/*"
  fields:
    apache: true
```

```
filebeat.inputs:
- type: filestream
...
parsers:
- ndjson:
  target: ""
  message_key: msg
```

# Filebeat



Filebeat.yml >> Input

## Filestream Input

<https://www.elastic.co/guide/en/beats/filebeat/current/filebeat-input-filestream.html>

```
filebeat.inputs:  
- type: filestream  
  ...  
  prospector.scanner.exclude_files: ['\\.gz$']  
  prospector.scanner.include_files: ['^/var/log/.*']
```

```
filebeat.inputs:  
- type: filestream  
  ...  
  include_lines: ['^ERR', '^WARN']  
  exclude_lines: ['^DBG']
```

# Filebeat



Filebeat.yml >> Input

## Filestream Input

<https://www.elastic.co/guide/en/beats/filebeat/current/multiline-examples.html>

### Problem

```
[beat-logstash-some-name-832-2015.11.28] IndexNotFoundException[no such index]
  at
org.elasticsearch.cluster.metadata.IndexNameExpressionResolver$WildcardExpressionResolver.resolve(IndexNameExpressionResolver.java:566)
  at org.elasticsearch.cluster.metadata.IndexNameExpressionResolver.concreteIndices(IndexNameExpressionResolver.java:133)
  at org.elasticsearch.cluster.metadata.IndexNameExpressionResolver.concreteIndices(IndexNameExpressionResolver.java:77)
  at org.elasticsearch.action.admin.indices.delete.TransportDeleteIndexAction.checkBlock(TransportDeleteIndexAction.java:75)
```

### Solution

```
parsers:
- multiline:
  type: pattern
  pattern: '^\\['
  negate: true
  match: after
```

# Filebeat



Filebeat.yml >> Processors

## Processors Types

<https://www.elastic.co/guide/en/beats/filebeat/current/defining-processors.html>

- add\_docker\_metadata
- add\_fields
- add\_host\_metadata
- add\_id
- add\_locale
- append
- convert
- copy\_fields
- decode\_base64\_field
- decode\_csv\_fields
- decode\_duration
- decode\_json\_fields
- decode\_xml
- decompress\_gzip\_field
- dissect
- drop\_event
- drop\_fields
- extract\_array
- fingerprint
- include\_fields
- move-fields
- rename
- replace
- script
- timestamp
- translate\_sid

# Filebeat



Filebeat.yml >> Processors

## Processors Conditions

<https://www.elastic.co/guide/en/beats/filebeat/current/defining-processors.html>

- equals
- contains
- regexp
- range
- network
- has\_fields
- or
- and
- not

# Filebeat



Filebeat.yml >> Processors

## Examples

```
processors:
  - add_fields:
      target: project
      fields:
        name: myproject
        id: '574734885120952459'
```

```
processors:
  - convert:
      fields:
        - {from: "src_ip", to: "source.ip", type: "ip"}
        - {from: "src_port", to: "source.port", type: "integer"}
      ignore_missing: true
      fail_on_error: false
```

# Filebeat



Filebeat.yml >> Processors

## Examples

```
processors:  
  - rename:  
      fields:  
        - from: "a.g"  
          to: "e.d"  
      ignore_missing: false  
      fail_on_error: true
```

```
  - replace:  
      fields:  
        - field: "file.path"  
          pattern: "/usr/"  
          replacement: "/usr/local/"  
      ignore_missing: false  
      fail_on_error: true
```

```
processors:  
  - drop_event:  
      when:  
        equals:  
          log_status: "INFO"
```

# Filebeat



Filebeat.yml >> Processors

## Examples

### Problem

```
"321 - App01 - WebServer is starting"
"321 - App01 - WebServer is up and running"
"321 - App01 - WebServer is scaling 2 pods"
"789 - App02 - Database is will be restarted in 5 minutes"
"789 - App02 - Database is up and running"
"789 - App02 - Database is refreshing tables"
```

### Solution

```
processors:
  - dissect:
      tokenizer: '"${service.pid|integer} - ${service.name} - ${service.status}"'
      field: "message"
      target_prefix: ""
```

# Filebeat



Filebeat.yml >> Output

## Output Types

- Elasticsearch Service
- Elasticsearch
- Logstash
- Kafka
- Redis
- File
- Console

# Filebeat



Filebeat.yml >> Output

Elasticsearch Output

```
output.elasticsearch:  
  hosts: ["https://localhost:9200"]  
  index: "your_index_name"  
  ssl.enabled: true # Enable SSL/TLS  
  ssl.certificateAuthorities: ["/path/to/your/ca.crt"]  
  username: "your_username" # Elasticsearch username  
  password: "your_password" # Elasticsearch password
```

Logstash Output

```
output.logstash:  
  hosts: ["localhost:5044"]
```

# Filebeat



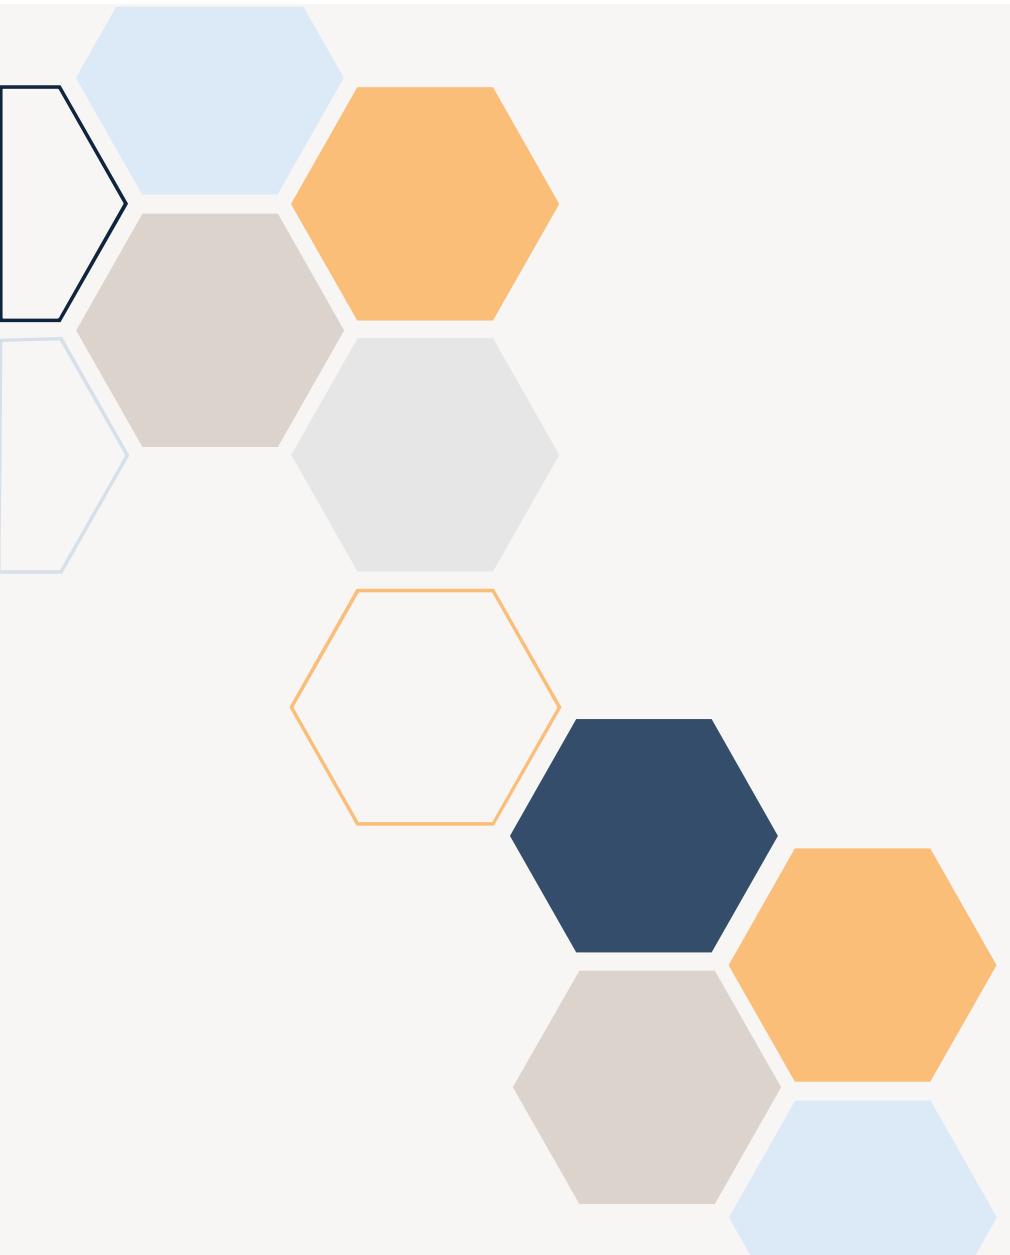
LAB

## Tasks

You have a json log file named [`spring-boot-app-logs.json`](#). Use filebeat to process the file and send it to your logstash instance listening on port 5044.

# Questions





# Thank you

Mr. Iyanou Eraste AKANDE  
[eraste.akande@gmail.com](mailto:eraste.akande@gmail.com)

# Starting with Logstash

Mr. Iyanou Eraste AKANDE  
Elastic Certified Engineer  
Data Engineer at Synaptique Maghreb



# Setup Logstash

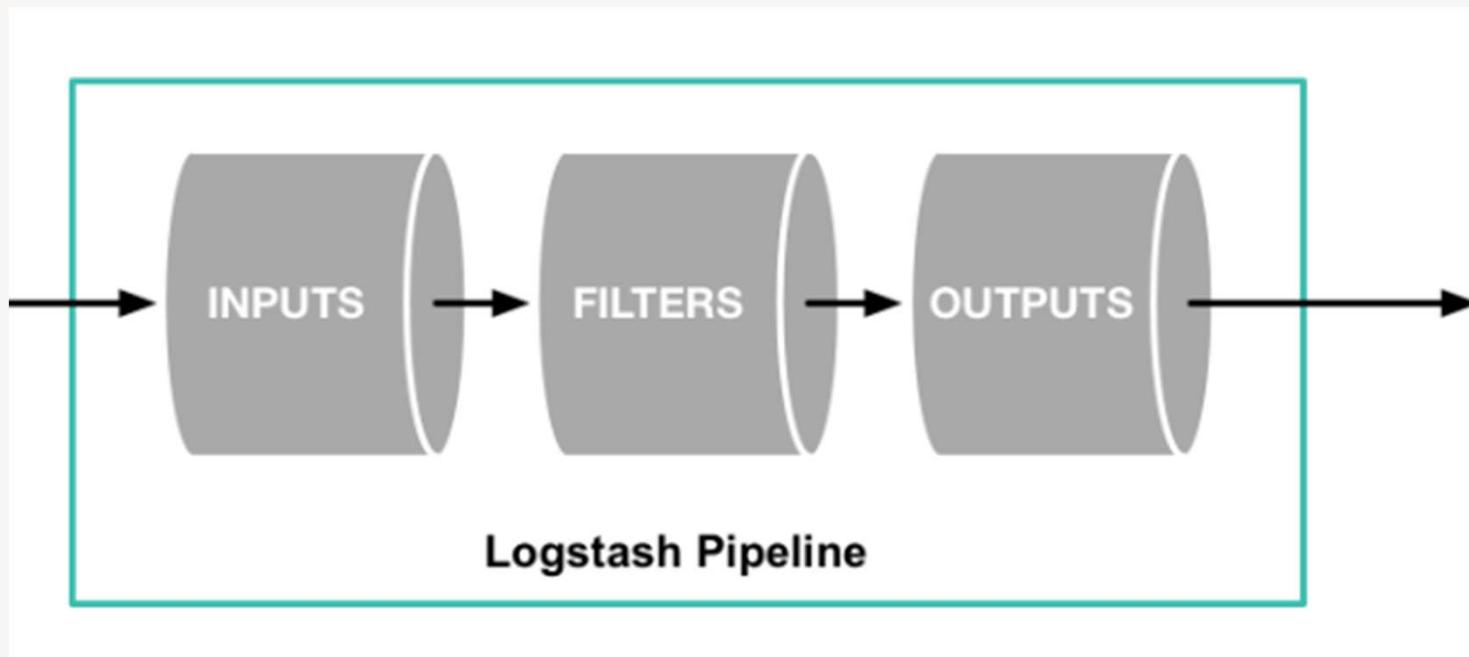


- Download Logstash and decompress it (`wget https://artifacts.elastic.co/downloads/logstash/logstash-8.12.2-linux-x86\_64.tar.gz && tar -xvf logstash-8.12.2-linux-x86_64.tar.gz`)
- Allow firewall on port 9600
- Start Logstash pipeline (“`cd logstash_directory`” && “`bin/logstash -f config/configuration-file.conf`”)

# Configuration



Architecture



# Configuration



Sample file

```
# Sample Logstash configuration for creating a simple
# Beats -> Logstash -> Elasticsearch pipeline.
input {
    beats { port => 5044 }
}
filter {}
output {
    elasticsearch {
        hosts => ["http://localhost:9200"]
        index => "filebeat-%{+YYYY.MM.dd}"
        user => "elastic"
        password => "changeme"
    }
}
```

# Configuration



## Input

```
input {  
    beats {  
        port => 5044  
    }  
}
```

<https://www.elastic.co/guide/en/logstash/current/input-plugins.html>

# Configuration



## Input

```
input {  
  
    file {  
        path => ["/path/samples/input-file.csv"]  
        start_position => "beginning"  
        sincedb_path => "/path/log.sincedb"  
        file_completed_action => "log"  
        file_completed_log_path => "/path/log.processed"  
        mode => read  
    }  
  
}
```

<https://www.elastic.co/guide/en/logstash/current/input-plugins.html>

# Configuration



## Filter

```
csv {  
    source => "message"  
    separator => ";"  
    columns => ["field1", "field2", "field3", "field4"]  
}  
  
mutate {  
    add_field => { "country" => "Morocco" }  
}  
  
mutate {  
    convert => { "money" => "float" }  
}  
  
date {  
    match => ["@timestamp", "yyyyMMdd HH:mm:ss"]  
    timezone => "UTC"  
    target => "@timestamp"  
}  
  
mutate {  
    rename => {"Lieu de Naissance" => "birth_city"}  
}
```

<https://www.elastic.co/guide/en/logstash/current/filter-plugins.html>

# Configuration



## Filter

```
mutate { remove_field => [ "message", "@version", "host", "path", "event"]}

if ([code] =~ /.+/ ) {
translate {
    source => "[code]"
    target => "[value]"
    dictionary_path => "/path/dicts/code.yml"
    exact => true
    override => true
    regex => false
}
}

translate {
    source => "[code]"
    target => "[code]"
    override => true
    exact => true
    dictionary => {
        "0" => "Normal"
        "1" => "Free"
    }
}
```

<https://www.elastic.co/guide/en/logstash/current/filter-plugins.html>

# Configuration



Filter

```
ruby {  
    code => "  
        event.set('salary', event.get('salary') / event.get('total_number'));  
        event.set('id', '2024' + event.get('id'));  
    "  
}
```

```
# location : "region=Rabat,country=Morocco"  
kv {  
    source => "location"  
    field_split => ","  
    value_split => "="  
}  
  
if [loglevel] == "debug" {  
    drop { }  
}
```

<https://www.elastic.co/guide/en/logstash/current/filter-plugins.html>

# Configuration

Filter



```
mutate {  
    add_field => { "hostname" => "LENOVO" }  
    #description : "The//name//of//the city"  
    gsub => [  
        "description", "//", " "  
    ]  
    uppercase => [ "fieldname" ]  
    lowercase => [ "fieldname" ]  
    copy => { "source_field" => "dest_field" }  
}  
  
# message : "55.3.244.1 GET /index.html 15824 0.043"  
grok {  
    match => { "message" => "%{IP:client} %{WORD:method} %{URIPATHPARAM:request} %{NUMBER:bytes} %{NUMBER:duration}" }  
}
```

<https://www.elastic.co/guide/en/logstash/current/filter-plugins.html>

# Configuration



## Output

```
elasticsearch {  
    hosts          => "server-address"  
    user           => "elastic"  
    password       => "changeme"  
    cacert         => "/path/cert/elk.crt"  
    ssl            => "true"  
    ssl_certificate_verification => "true"  
    index          => "logstash-test"  
    action         => "index"  
}
```

<https://www.elastic.co/guide/en/logstash/current/output-plugins.html>

# Configuration



Output

```
kafka {  
    codec => json  
    bootstrap_servers => "server_address:9092"  
    topic_id => "topic-name"  
    compression_type => "gzip"  
}
```

<https://www.elastic.co/guide/en/logstash/current/output-plugins.html>

# LAB



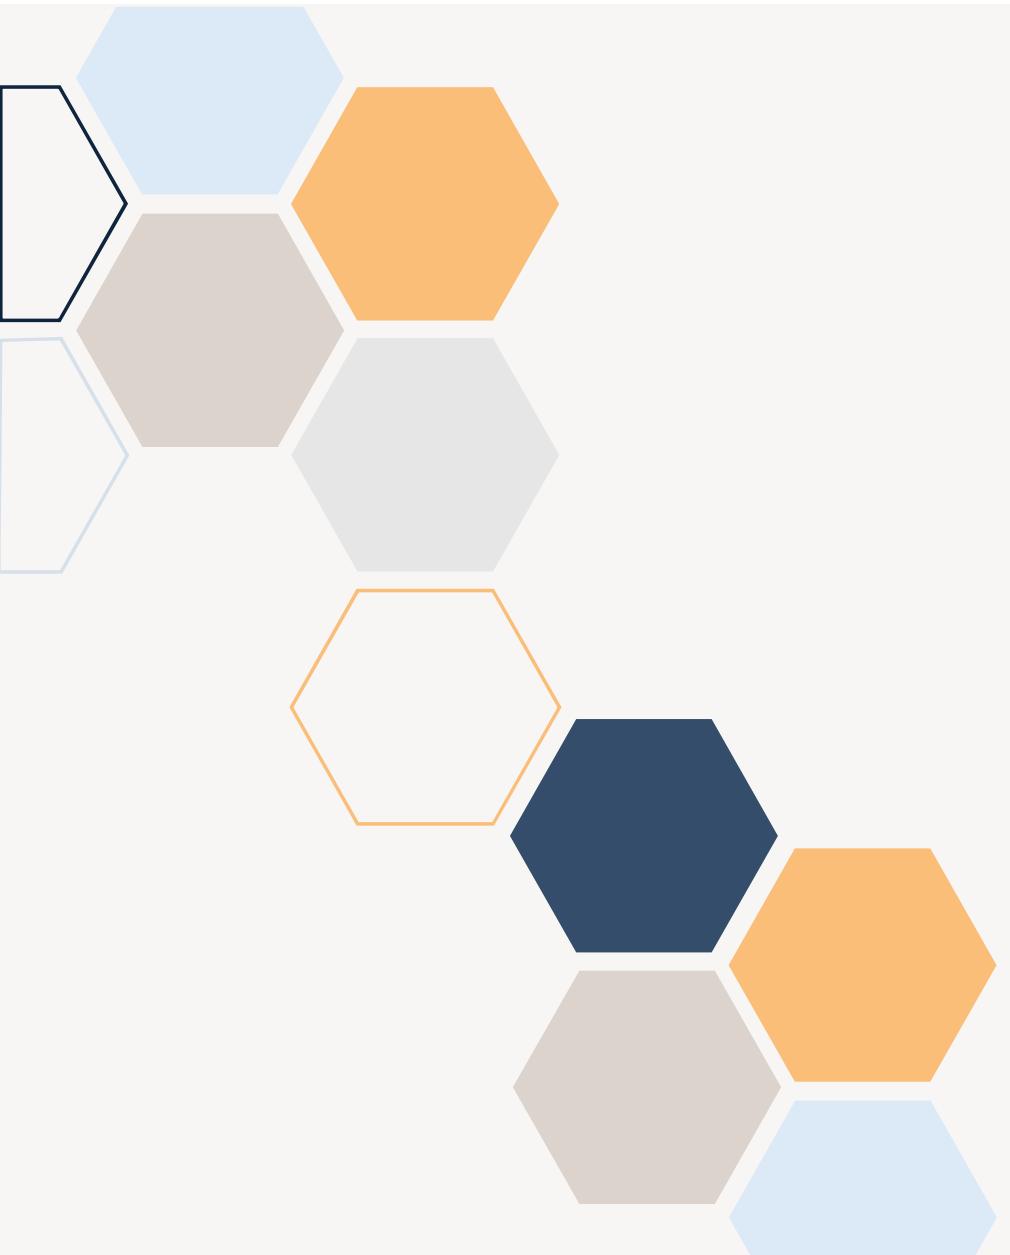
## Tasks

You are receiving a spring boot app logs from filebeat on port 5044. Configure logstash to process these logs and send the results to elasticsearch.

1. Use json filter to parse the message and add a field named **dataset** with the value “Spring App”
2. Rename the field **log\_level** to **level**
3. Remove the field **event** and format the field **event\_date** with the date filter
4. Apply a **gsub** filter to replace “com.example.” with a whitespace “” in the field **logger**
5. Define for the data stream “filebeat-8.13.0” an analyser named **filebeat\_analyser**. This analyser will use the **standard tokenizer** and the **lowercase** and **synonym token filters** so that “failed” and “error” are synonyms.
6. Define an index template for the data stream “filebeat-8.13.0”. Map these fields like this (logger → keyword, thread\_id → keyword, user\_id → keyword, request\_id → keyword, @timestamp → date, event\_date → date, message → text, level → keyword)
7. Apply the **filebeat\_analyser** to the field **message**
8. Send the logs to your Elasticsearch Output

# Questions





# Thank you

Mr. Iyanou Eraste AKANDE  
[eraste.akande@gmail.com](mailto:eraste.akande@gmail.com)

# Starting with Elasticsearch

Mr. Iyanou Eraste AKANDE  
Elastic Certified Engineer  
Data Engineer at Synaptique Maghreb



# Setup Elasticsearch



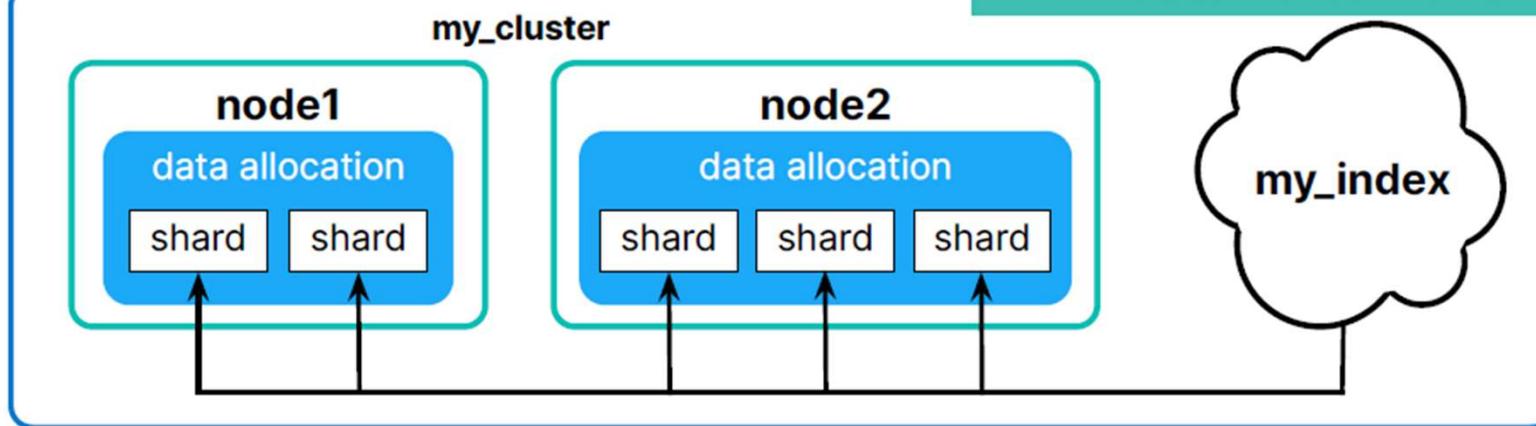
- Download Elasticsearch and decompress it (<https://www.elastic.co/guide/en/elasticsearch/reference/current/install-elasticsearch.html>)
- Make sure the user has write access on elasticsearch folders
- Allow firewall on port 9200
- Start Elasticsearch (“cd elasticsearch\_directory” && “bin/elasticsearch”)
- Copy username, password and Kibana enrollment token in a file
- Verify Elasticsearch Installation (<https://localhost:9200>, [https://localhost:9200/\\_cluster/health?pretty](https://localhost:9200/_cluster/health?pretty))
- Run Elasticsearch as Service
  - Windows (“bin/elasticsearch-service.bat install” && “bin/elasticsearch-service.bat start”)
  - Linux
    - Run as daemon (“bin/elasticsearch -d -p pid” to start and “pkill -F pid” to stop)
    - Use Systemctl (Define elasticsearch.service && “systemctl daemon-reload”, “systemctl enable elasticsearch.service”, “systemctl start elasticsearch.service”)

# Architecture



Shards

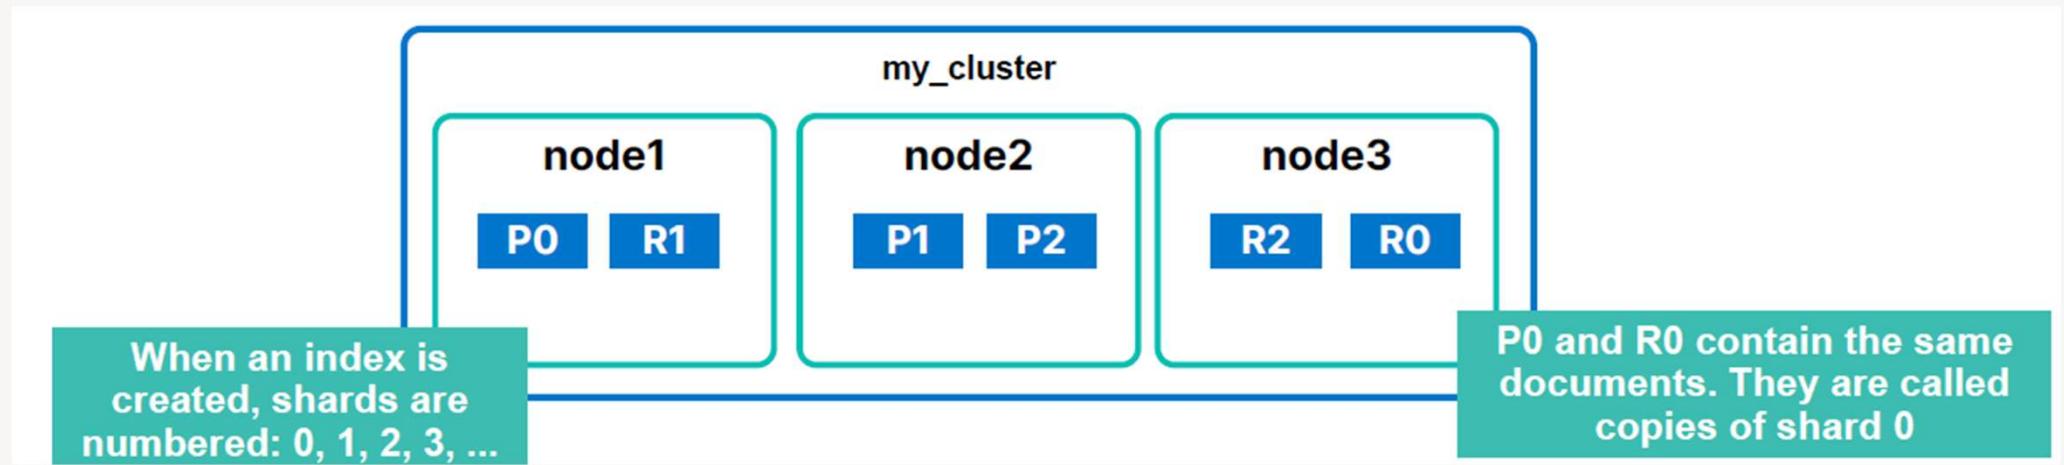
This index has 5 shards: every shard holds approximately 20% of the documents in this index



# Architecture



Primary Vs Replica



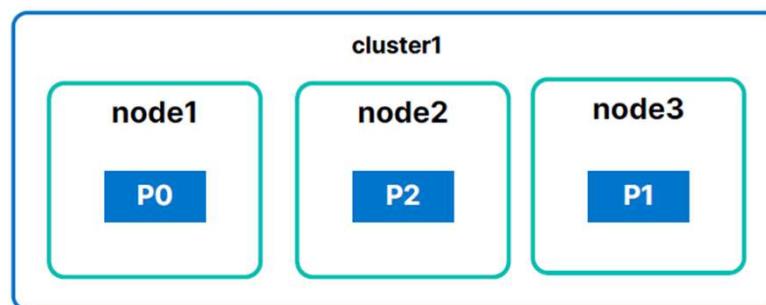
# Architecture



## Shards Configuration

request

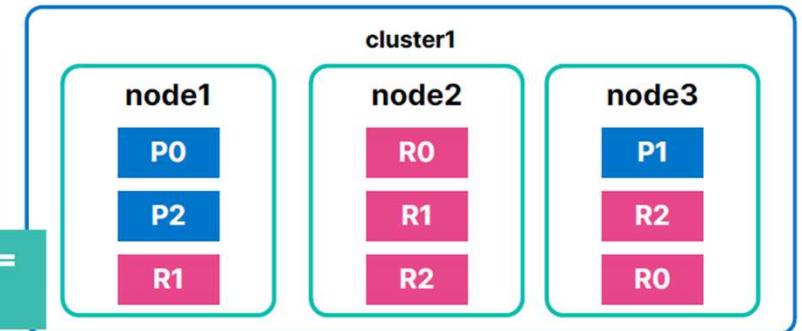
```
PUT my_new_index
{
  "settings": {
    "number_of_shards": 3
  }
}
```



request

```
PUT my_new_index/_settings
{
  "number_of_replicas": 2
}
```

3 primaries + 2 replica sets =  
9 total shards



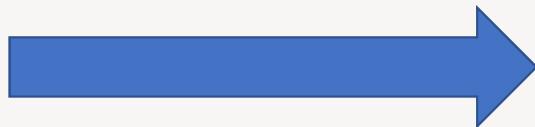
# Mapping

## Default Mapping

```
{  
  "first_name" : "Ali",  
  "name" : "KOZNI",  
  "gender" : "M",  
  "city" : "Rabat",  
  "country" : "Morocco",  
  "age" : 25,  
  "company" : "INTEL",  
  "starting_date" : "2022-04-23"  
}
```



- Slow index speed
- Non optimized storage space
- Not improved searches



```
"first_name" : text  
"first_name.keyword" : keyword  
"name" : text  
"name.keyword" : keyword  
"gender" : text  
"gender.keyword" : keyword  
"city" : text  
"city.keyword" : keyword  
"country" : text  
"country.keyword" : keyword  
"age" : long  
"company" : text  
"company.keyword" : keyword  
"starting_date" : text  
"starting_date.keyword" : keyword
```

Starting with Elasticsearch

# Mapping

## Explicit Mapping

```
PUT employees_data
{
  "mappings": {
    "properties": {
      "first_name": {
        "type": "keyword"
      },
      "age": {
        "type": "integer"
      }
      "starting_date": {
        "type": "date",
        "format": "yyyy-MM-dd"
      }
    }
  }
}
```

Starting with Elasticsearch



- Fast index speed
- Optimized storage space
- Improved searches



```
"first_name" : keyword
"name" : keyword
"gender" : keyword
"city" : keyword
"country" : keyword
"age" : integer
"company" : keyword
"starting_date" : date
```

# Mapping

## Data Types



- **text**: for full-text (analyzed) strings
- **keyword**: for exact value strings and aggregations
- **date** and **date\_nanos**: string formatted as dates, or numeric dates
- integer types: **byte**, **short**, **integer**, **long**
- floating-point numbers: **float**, **double**, **half\_float**
- **boolean**

Starting with Elasticsearch

# Mapping

## Properties



- Mappings are Elasticsearch's data schema
- Mappings are defined **per index**
- If you do not define an explicit mapping, Elasticsearch will **dynamically map** the fields in your documents
- You cannot change the mapping of a field after the index has been created, but you can add new fields to a mapping
- Creating a custom mapping will produce savings in search and index speed, as well as memory and storage requirements

Starting with Elasticsearch

# Component Template



- Reusable building blocks that can contain settings, mappings or aliases
- Reused across multiple templates

```
PUT _component_template/component_template1
{
  "template": {
    "mappings": {
      "properties": {
        "@timestamp": {
          "type": "date"
        }
      }
    }
  }
}
```



# Index Template

- Multiple indices with the same settings and mappings
- Index template can contain
  - **settings**
  - **mappings**
  - **aliases**
  - **component templates**

```
PUT _index_template/template_1
{
  "index_patterns": ["spring-*"],
  "template": {
    "settings": {
      "number_of_shards": 1
    },
    "mappings": {
      "properties": {
        "host_name": {
          "type": "keyword"
        }
      }
    },
    "priority": 500,
    "composed_of": ["component_template1"]
  }
}
```

# Analysers

How to analyse text fields ?



Example: A quick brown fox jumps over the lazy dog

- Character filters : Transform the stream by adding, removing, or changing characters (Ex: mapping, html\_strip)
- Tokenizers : Breaking a text down into smaller chunks (Ex: whitespace)
- Token filters : A *token filter* receives the token stream and may add, remove, or change tokens (Ex: lowercase, stop, synonym)

## Built in Analysers

<https://www.elastic.co/guide/en/elasticsearch/reference/current/analysis-analyzers.html>

## Built in Tokenizers

<https://www.elastic.co/guide/en/elasticsearch/reference/current/analysis-tokenizers.html>

## Built in Token filters

<https://www.elastic.co/guide/en/elasticsearch/reference/current/analysis-tokenfilters.html>

# Analysers

Create a custom analyser



```
PUT /my-index-000001
{
  "settings": {
    "analysis": {
      "analyzer": {
        "my_analyzer": {
          "tokenizer": "whitespace",
          "filter": [ "stop" ]
        }
      }
    }
  }
}
```

```
PUT trim_example
{
  "settings": {
    "analysis": {
      "analyzer": {
        "keyword_trim": {
          "tokenizer": "keyword",
          "filter": [ "trim" ]
        }
      }
    }
  }
}
```

# Analysers

Apply an analyser to a field



```
PUT my-index-000001
{
  "mappings": {
    "properties": {
      "title": {
        "type": "text",
        "analyzer": "whitespace"
      }
    }
  }
}
```

# Query DSL



## Search

```
GET index_name/_search
{
  "query": {
    "match": {
      "title": "lifestyle"
    }
  }
}
```

```
GET index_name/_search
{
  "query": {
    "term": {
      "name": "Alain"
    }
  }
}
```

<https://www.elastic.co/guide/en/elasticsearch/reference/current/query-dsl.html>

# Query DSL



## Aggregations

```
GET index_name/_search
{
  "aggs": {
    "my-agg-name": {
      "terms": {
        "field": "my-field"
      }
    }
  }
}
```

```
GET index_name/_search
{
  "aggs": {
    "my-second-agg-name": {
      "avg": {
        "field": "my-other-field"
      }
    }
  }
}
```

<https://www.elastic.co/guide/en/elasticsearch/reference/current/search-aggregations.html>



# Reindex API

Local

```
POST _reindex
{
  "source": {
    "index": " source_index_name",
    "query": {...}
  },
  "dest": {
    "index": " destination_index_name"
  }
}
```



# Reindex API

Remote

POST \_reindex

```
{  
  "source": {  
    "remote": {  
      "host": "http://otherhost:9200",  
      "username": "user",  
      "password": "pass"  
    },  
    "index": "remote_index",  
  },  
  "dest": {  
    "index": "local_index"  
  }  
}
```

Starting with Elasticsearch

18



# Update By Query

POST index\_name/\_update\_by\_query

```
{  
  "query": {  
    ....  
  },  
  "script": {  
    "source": "ctx._source.name = Ali"  
  }  
}
```



# Delete By Query

`POST index_name/_delete_by_query`

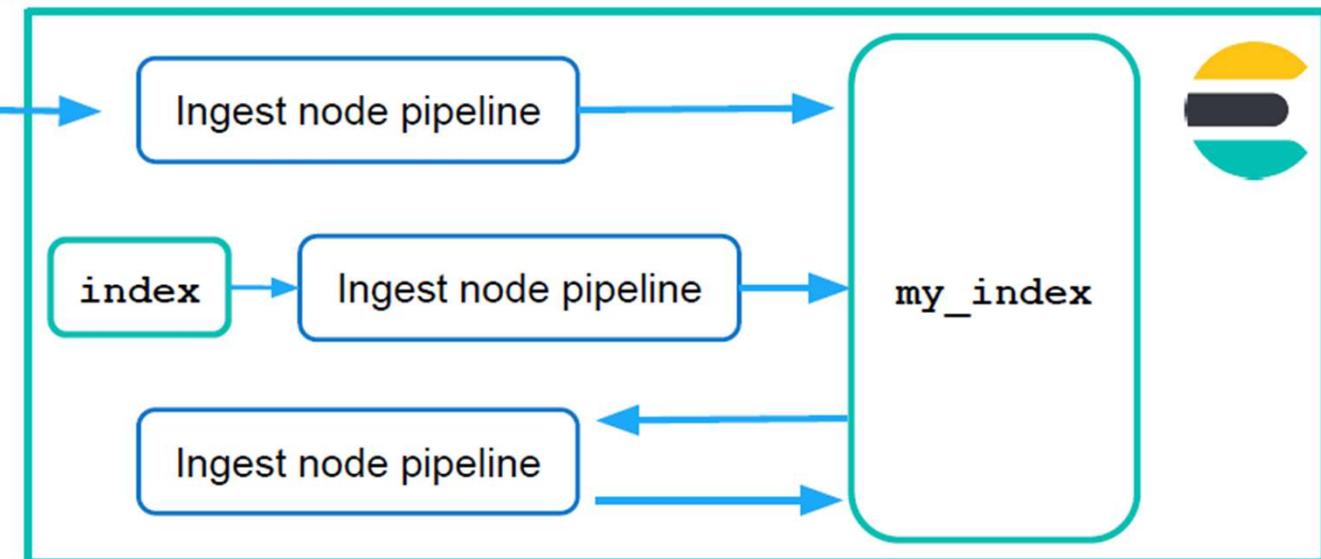
```
{  
  "query": {  
    "match": {  
      "name": "Malton Jack"  
    }  
  }  
}
```

# Ingest Node Pipelines



How it works ?

**There are many ways you can use an ingest node pipeline**



# Ingest Node Pipelines



## Pipeline Creation

Screenshot of the Elasticsearch "Create pipeline" interface:

**Management** sidebar:

- Ingest** (selected)
- Ingest Pipelines** (selected)
- Logstash Pipelines

**Data**

- Index Management
- Index Lifecycle Policies
- Snapshot and Restore
- Rollup Jobs
- Transforms
- Cross-Cluster Replication
- Remote Clusters

**Alerts and Insights**

- Rules and Connectors
- Cases
- Reporting
- Machine Learning
- Watcher

**Security**

- Users
- Roles
- API keys

**Create pipeline** form:

**Name**: A unique identifier for this pipeline.

Add version number

**Description**: A description of what this pipeline does.

**Processors** section:

**Add your first processor**

Use processors to transform data before indexing. [Learn more](#).

**Add a processor**

**Import processors**

**Create pipeline** button

# Ingest Node Pipelines



Processors

## Manipulate Fields

- set
- remove
- rename
- dot\_expander
- etc

## Manipulate Values

- split/join
- grok
- dissect
- gsub
- etc

## Special Operations

- csv/json
- geoip
- user\_agent
- script
- pipeline
- etc

# Ingest Node Pipelines



## Usage

- **POST** my\_index/\_update\_by\_query? **pipeline=**my\_pipeline

- **PUT** my\_index

```
{  
  "settings": {  
    "default_pipeline": "my_pipeline"  
  }  
}
```

- **POST** \_reindex

```
{  
  "source": {  
    "index": "my_index"  
  },  
  "dest": {  
    "index": "new_index",  
    "pipeline": "my_pipeline"  
  }  
}
```



# Enrich Data

Explanation

stock			provenance		
id	product	price	id	origin	
1	rice	40	1	Benin	
2	bag	30	2	Togo	
3	hat	15	3	Morocco	
4	dog	200	4	Congo	
5	desk	300	5	Mali	
6	chair	150	6	Niger	
7	table	250	7	Soudan	
8	computer	1000	8	Angola	
9	phone	750	9	Namibie	
10	oil	60	10	Rwanda	

# Enrich Data



## Enrich Policy

```
PUT _enrich/policy/ product_mapper
```

```
{  
  "match": {  
    "indices": "provenance",  
    "match_field": "id",  
    "enrich_fields": ["country"]  
  }  
}
```

# Enrich Data



Create Enrich Index

**POST \_enrich/policy/product\_mapper/\_execute**

# Enrich Data



## Ingest Pipeline

```
PUT /_ingest/pipeline/product_mapper_pipeline
{
  "processors" : [
    {
      "enrich" : {
        "policy_name": "product_mapper",
        "field" : "id",
        "target_field": "country"
      }
    }
  ]
}
```

# Enrich Data



Use the pipeline

**POST stock/\_update\_by\_query?pipeline=product\_mapper\_pipeline**

# Transforms

Explanation



# Transforms



Types

## Create transform

### 1 Configuration



#### Pivot

Aggregate and group your data

✓ Selected



#### Latest

Keep track of your most recent data

Select

# Transforms



Mode

Status ▾ Mode ▾

Reload

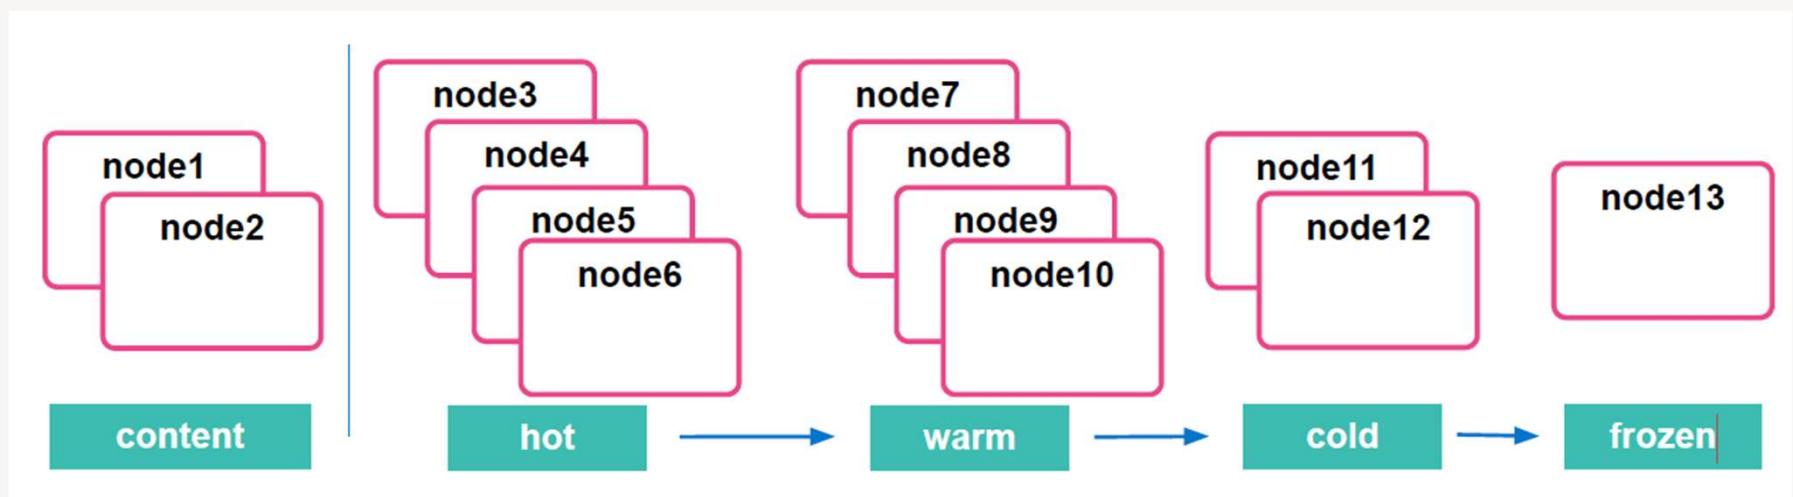
batch

continuous

# Index Lifecycle Management



Data Tiers



`node.roles: ["data_warm", "data_content"]`

# Security

## Role



### Create role

Set privileges on your Elasticsearch data and control access to your Kibana spaces.

Role name



Elasticsearch hide

#### Cluster privileges

Manage the actions this role can perform against your cluster. [Learn more](#)

Add an action...



#### Run As privileges

Allow requests to be submitted on the behalf of other users. [Learn more](#)

Add a user...



#### Index privileges

Control access to the data in your cluster. [Learn more](#)

# Security

User



## Create user

### Profile

Provide personal details.

Username



Full name

Email address

### Password

Protect your data with a strong password.

Password



Password must be at least 6 characters.

Confirm password



### Privileges

Assign roles to manage access and permissions.

Roles

▼

[Learn what privileges individual roles grant.](#)

# Snapshots & Restore



## Repository

**Register repository**

**Repository name**  
A unique name for the repository.

**Repository type**  
Storage location for your snapshots. [Learn more about repository types.](#)

**Azure** [Learn more](#) [Select](#)

**Google Cloud Storage** [Learn more](#) [Select](#)

**AWS S3** [Learn more](#) [Select](#)

**Shared file system** [Learn more](#) [Select](#)

**Read-only URL** [Learn more](#) [Select](#)

# Snapshots & Restore



## Policy

**Policy name**  
A unique identifier for this policy.

**Snapshot name**  
The name for the snapshots. A unique identifier is automatically added to each name.

**Repository**  
The repository where you want to store the snapshots.

**Schedule**  
The frequency at which to take the snapshots.

**Name**  
daily-snapshots

**Snapshot name**  
<daily-snap-{now/d}>  
Supports date math expressions. [Learn more.](#)

**Repository**  
⚠ You don't have any repositories  
You must register a repository to store your snapshots.  
⊕ Register a repository

**Frequency**  
Every day

**Time**  
At 01 : 30

[Create cron expression](#)

**Next >** **Cancel**

# Snapshots & Restore



Restore

## Snapshot and Restore

[Snapshot and Restore docs](#)

Use repositories to store and recover backups of your Elasticsearch indices and clusters.

[Schemas](#)   [Repositories](#)   [Policies](#)   [Restore Status](#)

Search...

Repository

Reload

<input type="checkbox"/> Snapshot	Repository	Indices	Shards	Failed shards	Date created	Duration	Actions
<input type="checkbox"/> apps-indices-2024.03.16- -Orwoim8t4kqrzlyoaplbg	NFS_Apps_Indices	16	22	0	Mar 16, 2024 1:29 PM GMT+1	3s	 



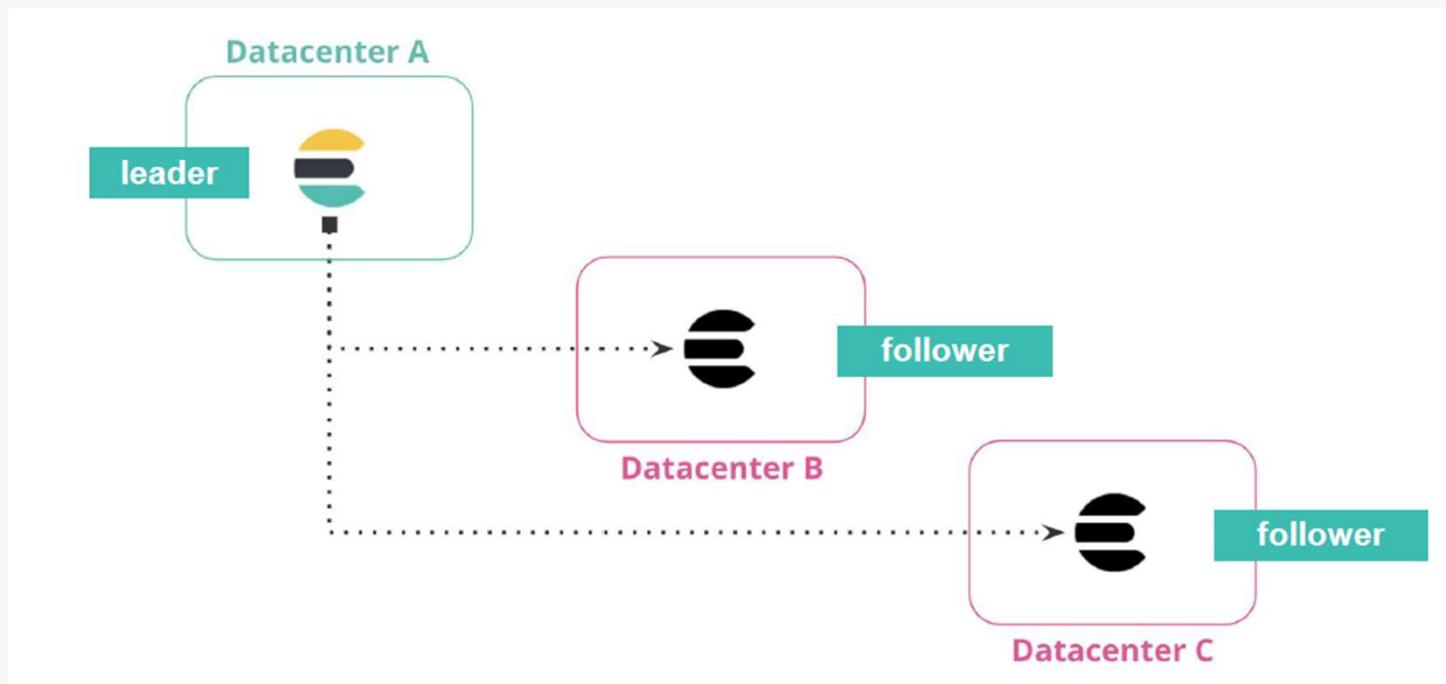
# Troubleshooting

Command	Response
GET _tasks	the running tasks on the cluster
GET _cluster/pending_tasks	any cluster-level changes that have not yet executed
GET _nodes/hot_threads	threads using high CPU volume and executing for a long time
GET _cat/shards	the statistics of shards
GET _cat/health	the health of the cluster
GET _cluster/allocation/explain	explain the shards allocation



# Cross Cluster Replication

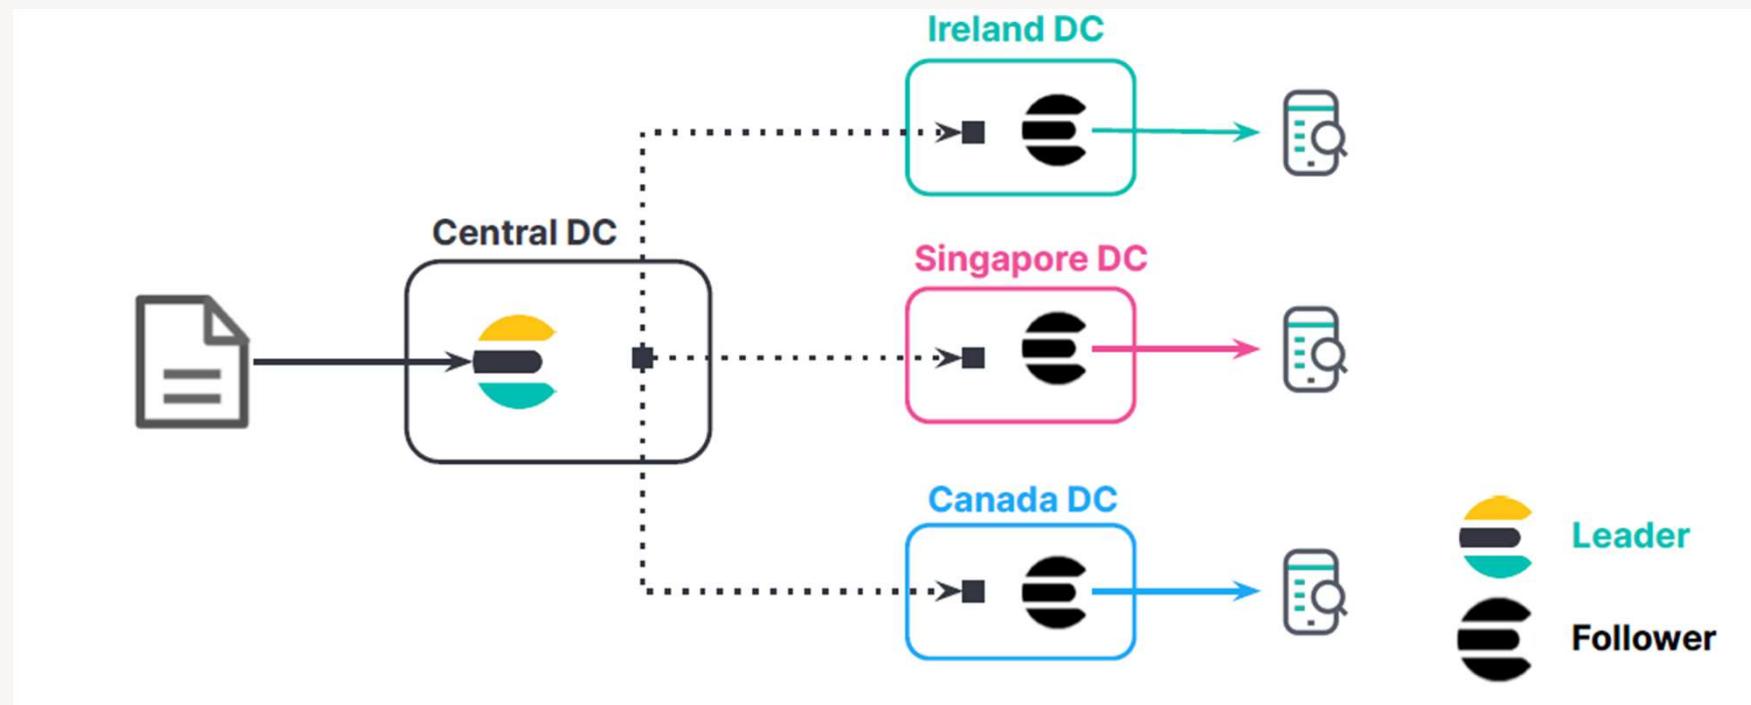
Disaster recovery and high availability



# Cross Cluster Replication



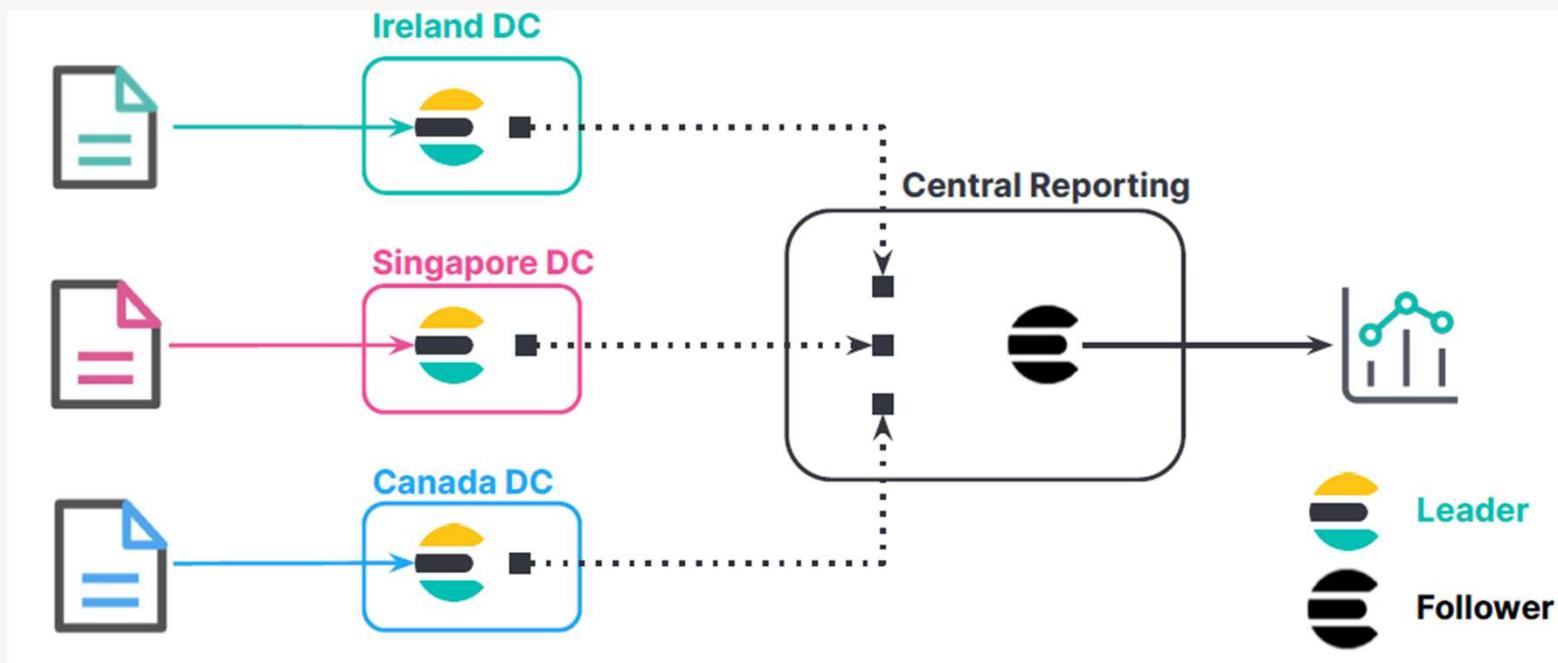
Data locality



# Cross Cluster Replication

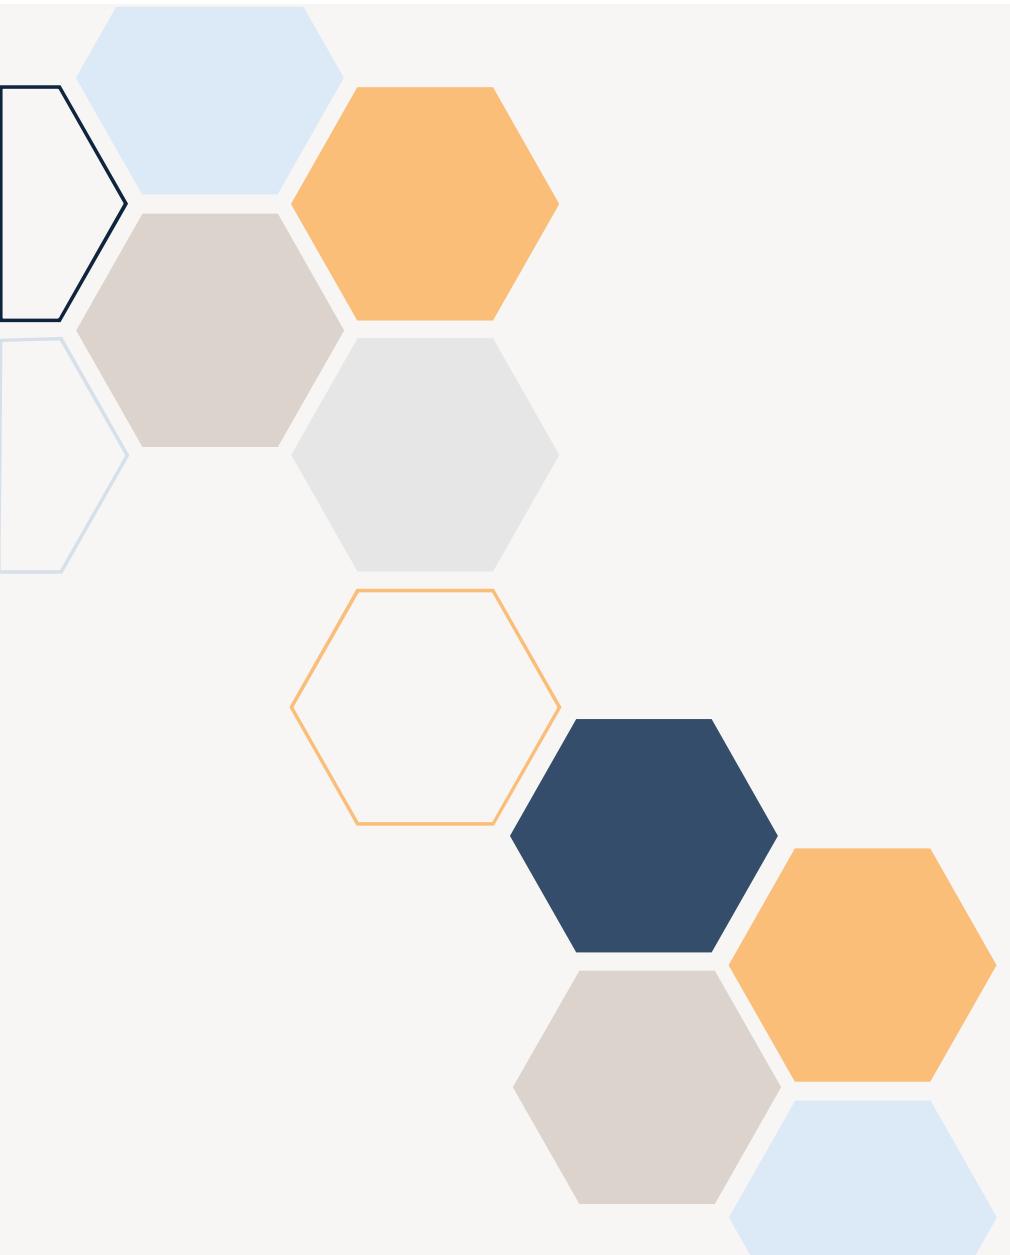


Centralized reporting



# Questions





# Thank you

Mr. Iyanou Eraste AKANDE  
[eraste.akande@gmail.com](mailto:eraste.akande@gmail.com)